

talabat Group

AI-Enabled Investment Allocation ·

Version 1.1

How to allocate and govern the 2026 USD175 million programme across Everyday App and Food-leadership initiatives to maximize profitable growth and long-term platform economics.

Programme	AASTMT MBA · AI for Business Organizations
Team	Group G02
Scope	talabat Group · eight operating markets
Decision horizon	2026 investment programme
Document status	Version 1.1 enhancement · independently verified baseline · 25 July 2026

Recommendation in one line

Protect Food leadership in the evidenced GCC-3 markets, pilot talabat pro in Egypt then Iraq, continue proven Everyday App investments, and release funding through human-approved KPI gates.

Contents

1	Executive Summary
2	Business Description
3	Market Analysis
4	Value Proposition
5	AI Technology and Development
6	Business Model and Revenue Streams
7	Marketing and Sales Strategy
8	Operations Plan
9	Financial Plan
10	Risk Analysis
11	Corporate Social Responsibility & Responsible AI
12	Implementation Plan (Three Horizons)
13	Monitoring and Evaluation
14	Appendices

Evidence convention: disclosed facts are cited by TLB source/page or vault note; assumptions and forecasts retain ASM/node identifiers; Agentic OS allocations are explicitly labeled analytical.

1. Executive Summary


Approximately 69% of the disclosed 2026 programme is directed to Everyday App expansion



Source: TLB-020 pp.12,16; TLB-014 pp.16,19 | Classification: Disclosed talabat programme structure
Management implication: Govern release within the disclosed envelope; do not present the split as an Agentic OS allocation.

Exhibit 1. Approximately 69% of the disclosed 2026 programme is directed to Everyday App expansion.

The portfolio requires four actions—not one uniform funding decision



Source: Investment Options Register; DEC-009; OPT-001-005 | Classification: Agentic OS analytical recommendation; no dollar allocation shown
Management implication: Fund OPT-003 now, pilot OPT-002, continue OPT-001/005 at pace, and build OPT-004 later.

Exhibit 2. The portfolio requires four actions—not one uniform funding decision.

Situation

talabat Holding plc is a publicly listed, Group-wide on-demand delivery platform (Dubai Financial Market, IPO priced 29 November 2024) operating across eight markets — UAE, Kuwait, Qatar, Bahrain, Oman, Jordan, Iraq, and Egypt (Group; Problem_Charter.md). FY2025 results show a large, growing, and newly profitable business at scale: GMV of USD 9.5 billion (+28% cFX), revenue of USD 3.9 billion, Adjusted EBITDA of USD 615 million (6.5% margin), and 7.5 million active customers (Group; Problem_Charter.md). Management itself frames the company's strategic trajectory as a deliberate transition "from a multi-vertical food-delivery platform to the region's Everyday App" (Group; TLB-020, page 16, cited in Section_02_Business_Description.md §2.1) — the strategic intent this plan's recommendation operates inside, not one it invents. For 2026, the Board approved a USD 175 million investment programme, split by talabat's own disclosure into an Everyday App bucket (~USD 120 million: talabat mart dark-store densification, talabat pro loyalty scaling, new verticals) and a Food-leadership bucket (~USD 55 million: consumer value proposition, partner retention, targeted incentives), fully funded from internal cash (Group; vault/Decisions/Investment_Portfolio_Register.md; TLB-020, page 16).

Complication

talabat has not disclosed — and this project cannot observe — how that USD 175 million should be allocated within each bucket, across eight markets that differ sharply in maturity, profitability, and competitive intensity. The evidence assembled across this plan shows real dispersion of return, not a uniform growth story. Group Adjusted EBITDA margin compressed to 4.8% of GMV in Q1 2026 — down from 6.3% a year earlier and from 6.0% in Q4 2025 — a real, already-realized step-down, not a hypothetical risk (Group; Problem_Charter.md; Section_10_Risk_Analysis.md §10.3, FIN-01). Food-leadership's non-high-value customer retention has already eroded 4% year-on-year in three of talabat's most mature markets — UAE, Kuwait, and Qatar specifically (market-comparison; TLB-019, page 9; Section_03_Market_Analysis.md). Grocery & retail — the fastest-growing vertical and the largest named Everyday App line item — remains penetrated at only 1-3% of a USD 104-150 billion+ addressable category (Group/external; Section_03_Market_Analysis.md §3.1). And talabat discloses no internal capital-allocation committee, approval threshold, or stage-gate process for this or any programme (Group; Topics/Capital Allocation and Investment Governance.md; Section_08_Operations_Plan.md §8.4) — meaning a USD 175 million programme is currently being spent with no disclosed mechanism to redirect it if early evidence diverges from plan. Left undifferentiated, this programme risks funding what is already familiar rather than what the disclosed evidence actually supports.

Question

How should talabat allocate its 2026 USD 175 million investment programme across Everyday App and Food-leadership initiatives — and within them, across markets and initiative categories — to maximise profitable growth, customer lifetime value, customer retention, and long-term platform economics, given the evidence actually available, and without pretending to possess internal customer-level data this project does not have? (Problem_Charter.md)

Answer

This plan's answer is a framework and a committed sequence, not a single number. talabat does not need a new capital-allocation process invented from scratch — it needs the evidence-ranked, stage-gated discipline applied to the process it has already disclosed and already funded. This plan's governing hypothesis (Section_02_Business_Description.md §2.3) — that applying such a framework to the already-approved USD 175 million envelope will raise its realized GMV, EBITDA, and customer-lifetime-value contribution relative to an unranked allocation — is supported, not merely asserted, by every section that follows it: five evidence-grounded candidate initiatives were identified, compared on an explicit 11-criterion basis, and ranked (OPT-001-OPT-005, DEC-009, approved 2026-07-23, middle path).

Business Name and Overview

talabat Holding plc — a mature, profitable, publicly listed Group operating a three-sided marketplace (customers, restaurant/retail Partners, riders) across eight countries, generating GMV of USD 9.5 billion in FY2025 across food delivery, grocery and quick-commerce (talabat mart), and a growing set of adjacent services (subscription, advertising, FinTech) (Group; Section_02_Business_Description.md §2.1).

Mission and Vision

talabat's own stated trajectory — moving "from a multi-vertical food-delivery platform to the region's Everyday App," a framing management itself calls "a deliberate strategic choice, made from a position of strength" (Group; TLB-020, page 16) — is the mission and vision this plan's recommendation is built to serve, not a new one this plan proposes. This plan's own contribution sits one layer below that mission: making sure the USD 175 million funding that mission for 2026 is spent against explicit evidence and criteria, not left undifferentiated.

AI Product/Service, Target Market, and Value Proposition

The AI product is not customer-facing — it is an AI-enabled decision-support system that assembles talabat's own disclosed evidence, ranks candidate investment initiatives against explicit, documented criteria, builds range-bound (never single-point) financial scenarios, and monitors a KPI and stage-gate framework, gated for human approval before any capital moves (Problem_Charter.md; Section_02_Business_Description.md §2.4-2.5; Section_05_AI_Technology_and_Development.md §5.1). The target market is internal: talabat's own 2026 capital-allocation decision, applied across its eight operating markets, with Egypt as this plan's richest single-country worked example, not its boundary. The value proposition is comparative in kind, not degree: talabat's own disclosed rationale for the programme is qualitative ("investments rather than costs... expected to offset the margin impact over time," TLB-020, page 16), with no internal capital-allocation methodology disclosed anywhere in the corpus — this plan's USP is a transparent, evidence-traceable allocation framework where every recommended dollar resolves to a named Fact, Decision, or labeled Assumption, against that undisclosed status quo (Section_04_Value_Proposition.md §4.1). The customer-facing capability the 2026 programme itself scales — AI-driven personalisation, offer timing, and dispatch optimisation — is already built, licensed from parent Delivery Hero SE, and generating a disclosed, growing EBITDA contribution (>USD 14 million FY2024 → >USD 30 million FY2025); this plan does not need to build it, only to help allocate capital toward scaling it where the evidence is strongest (Section_05_AI_Technology_and_Development.md §5.1-5.2).

Key Objectives, Growth Goals, and Financial/Funding Summary

Objective: raise the realized GMV, EBITDA, and customer-lifetime-value contribution of the already-approved USD 175 million programme by replacing undifferentiated spend with an evidence-ranked, stage-gated sequence, without requesting incremental capital beyond what talabat has already committed (Section_02_Business_Description.md §2.3; Section_09_Financial_Plan.md §9.6). Growth goals: talabat's own disclosed FY2026 guidance — GMV +11-14% cFX, revenue +14-17% cFX — stated here as the Group base case this plan adopts, not a target this plan sets independently (Section_09_Financial_Plan.md §9.3). Financial and funding summary: the USD 175 million programme (~USD 120 million Everyday App, ~USD 55 million Food-leadership) is fully funded from internal cash; this plan requests no incremental capital and constructs no return figure beyond talabat's own disclosed FY2026 Adjusted EBITDA guidance of 4.4-4.8% of GMV and free cash flow guidance of 3.2-3.6% of GMV. Consistent with DEC-009's mandatory placement rule, no dollar range for any individual candidate initiative appears in this summary or anywhere else in this plan as a headline commitment — the five initiatives' illustrative cost ranges exist only as a labeled sensitivity exhibit in Section 14, each carrying the sentence "This is an Agentica OS analytical recommendation, not a disclosed talabat allocation"

Traceability (full — this section synthesizes, rather than newly cites, Sections 2-14)

2. Business Description

The Agentic OS governs capital decisions while talabat's embedded AI executes operating use cases




Exhibit 3. The Agentic OS governs capital decisions while talabat's embedded AI executes operating use cases.

Answer, stated first

talabat does not need to invent a new capital-allocation process — its Board has already approved and disclosed one: a USD175 million 2026 investment programme split into an Everyday App bucket (~USD120 million) and a Food-leadership bucket (~USD55 million), funded entirely from internal cash and framed explicitly as "investments rather than costs" (TLB-020, page 16). What talabat has not disclosed is how that programme is being allocated below the bucket level — across specific initiatives, across its eight operating markets, or against a stated evidence-and-return framework. This plan's governing hypothesis is that applying a disciplined, evidence-ranked, stage-gated allocation and monitoring framework to this already-approved envelope — rather than treating it as a single undifferentiated spend line — will raise the realized GMV, EBITDA, and customer-lifetime-value contribution of the programme relative to an unranked allocation, because the programme's own disclosed structure already separates a quantified cost (the FY2026 Adjusted EBITDA margin bridge) from an entirely undisclosed return, and because the candidate initiatives within it show materially different evidence strength, scalability, and reversibility profiles that an unranked allocation would not distinguish.

2.1 Industry Overview and Opportunities for AI Adoption

talabat operates a three-sided marketplace across eight countries (UAE, Kuwait, Qatar, Bahrain, Oman, Jordan, Iraq, Egypt — Entities/Countries.md), generating GMV of USD 7,428mn in FY2024 (+23% YoY) across food delivery, grocery and quick-commerce (talabat mart), and a growing set of adjacent services (subscription, advertising, FinTech) (Topics/GMV.md; Topics/2026 Investment Programme.md). Since its Dubai Financial Market listing (November 2024, TLB-025) management has repeatedly framed the company's strategic trajectory as a transition "from a multi-vertical food-delivery platform to the region's Everyday App" (TLB-020, page 16) — a framing management itself calls "a deliberate strategic choice, made from a position of strength." The 2026 investment programme is the financial expression of that choice: a Board-approved, internally-funded USD175 million commitment, first disclosed February 2026 (TLB-019) and fully decomposed by May 2026 (TLB-020, TLB-014) into ~USD120mn for Everyday App capability-building (talabat mart densification, talabat pro scaling, new verticals) and ~USD55mn for defending Food-leadership against "both new entrants and incumbents" (Topics/2026 Investment Programme.md).

2.2 Problem Statement (MECE)

The business problem this plan addresses decomposes into four mutually exclusive, collectively exhaustive sub-problems, each corresponding to a genuine, separately-evidenced gap in what talabat discloses:

- Cross-bucket allocation (already resolved by talabat, not this plan's decision). The USD175mn
- Within-bucket, cross-initiative allocation (open — this plan's primary contribution). Neither

2.3 Governing Hypothesis

We believe that applying an AI-enabled, evidence-ranked, stage-gated allocation and monitoring framework to talabat's disclosed USD175 million 2026 investment programme — rather than leaving its within-bucket and cross-market allocation undifferentiated — will raise the programme's realized GMV, EBITDA, and customer-lifetime-value

contribution, because (a) the programme's own disclosure already separates a quantified cost (the FY2026 Adjusted EBITDA margin bridge, 6.0%→4.6% of GMV guided) from an entirely undisclosed return, meaning no evidence currently exists to confirm the programme is being spent optimally even on its own terms, and (b) the five identified candidate initiatives show materially different confidence, scalability, and reversibility profiles (DEC-009's 11-criterion comparison) that an unranked, undifferentiated allocation would not distinguish or act on.

2.4 Core Features and Benefits

The proposed AI-enabled decision-support system's core features, each already evidenced by a real artifact this analysis produced, not a hypothetical capability:

- Evidence assembly and geography discipline — every claim tagged Group/GCC/non-GCC/Egypt-standalone/
- Investment Option comparison — a structured, 11-criterion, High/Medium/Low comparison framework

2.5 Business Model

The AI-enabled decision-support system is proposed as an internal capability, not a licensed or externally-sold product — consistent with talabat's own disclosed pattern of embedding AI/ML capability across functions rather than treating it as a discrete product line. talabat's own AI/ML capability is itself inherited from parent Delivery Hero SE's shared technology stack rather than built standalone (Topics/AI.md; Entities/Technology_Platforms.md) — this plan's proposed decision-support layer follows the same embedded-capability pattern, not the same build-vs-partner arrangement (see Section 5.2 for that distinction). vault/Forecasts/Value_Driver_Tree_v2.md N-06 notes AI/personalisation is "not named as a discrete dollar line within either bucket... an embedded capability." Consistent with that pattern, the capital-allocation decision-support layer proposed here is structured the same way: embedded into how the USD175mn programme (and any future programme like it) gets planned, monitored, and adjusted — not a standalone software product with its own revenue line. This is a deliberate, evidenced framing choice, not an oversight: Section 6 (Business Model and Revenue Streams) accordingly does not need to construct a new revenue stream for "the AI system" itself; its value is realized entirely through the better-allocated USD175mn it supports.


2.6 Current Business Stage

The underlying business (talabat) is public, mature, and profitable at Group level (IPO November 2024; FY2025 Adjusted EBITDA margin ~6.5% of GMV) — not a startup. The specific capability this plan proposes — an evidence-ranked, stage-gated capital-allocation framework applied to the 2026 investment programme — is at an early, pilot-appropriate stage: five candidate initiatives have been identified and ranked (DEC-009), but zero have yet moved past status: candidate to approved or scaled, and the cross-market allocation question (problem 3 in §2.2) remains open pending better evidence than the public corpus alone provides. This is the correct, honest stage to be at for a plan whose own charter states that "exact optimal allocations cannot be proven from public data alone" (Problem_Charter.md) — Section 12 proposes exactly this staged approach (pilot before scale, evidence-gated transitions) rather than a premature full-scale rollout recommendation.

Traceability (partial — full table due at Section 14)

3. Market Analysis

Non-GCC growth is more than twice GCC growth, but GCC still carries about four-fifths of GMV



Source: TLB-019 p.14; TLB-011 p.2 | Classification: Disclosed, FY2025 growth/share shown separately
Management implication: Protect GCC's scale while using reversible pilots to learn in faster-growing non-GCC markets.

Exhibit 4. Non-GCC growth is more than twice GCC growth, but GCC still carries about four-fifths of GMV.

Answer, stated first

talabat's own disclosures show a company that already dominates the categories it operates in by category-share multiple, yet has penetrated only a fraction of the underlying addressable demand — and the 2026 investment programme's existing shape already reflects that gap more than it reflects the scale advantage. Reconciling top-down category penetration against talabat's own realized bottom-up scale shows the largest unrealized headroom sits in grocery/retail (penetrated at low single digits of a USD 104–150bn+ addressable category) and in the faster-growing but still-minority non-GCC segment (+57% GMV growth y/y against GCC's +22% in FY2025, non-GCC rising from 18% of Group GMV in FY2025 to 21% of Group GMV by Q1 2026 — two dated data points, not a blended range) — precisely the two areas the ~USD 120mn Everyday App bucket is weighted toward. At the same time, talabat's own evidence shows category leadership alone has already failed to prevent a measured retention decline among non-high-value customers in three of its most mature markets, which is the direct, disclosed rationale for the ~USD 55mn Food-leadership bucket. Three of the plan's five ranked candidate Investment Options (OPT-001, OPT-003, OPT-004) map onto opportunity/threat pairs this section identifies directly; this section's genuinely open gaps — Egypt's inconsistent category-share figure, the absence of any Egypt-specific named competitor in the primary corpus, and the GCC-3-market scope of the Food-leadership competitive-pressure evidence — are stated explicitly below rather than smoothed over, consistent with Problem_Charter.md's own evidence-limitations standard.

3.1 Industry Trends — Market Sizing, Top-Down and Bottom-Up, Reconciled

Reframing note. The GSB template's Section 3 asks for "AI market size and growth projections" and "AI adoption trends... in the MENA region." Consistent with Section 2's framing (the AI-enabled capability this plan proposes is a capital-allocation decision-support layer, not a customer-facing product with its own addressable market), the market that matters for this decision is the market the USD 175mn programme actually operates in — talabat's own eight-country delivery, grocery, and retail marketplace — not a standalone "MENA AI market" figure. No AI-market-size or AI-adoption-by-sector TAM figure exists anywhere in the 29-document primary corpus; constructing one would require exactly the kind of fresh external research Part A's Stage 5 leaves open (Business_Plan_Generation_Pipeline.md). This section does not fabricate one — it sizes talabat's actual operating market instead, which is the evidence the corpus actually supports and the market the investment decision is about.

3.2 Target Market — MECE Segmentation

Primary segmentation axis: talabat's own disclosed reportable-segment geography — GCC / non-GCC (Jordan + Iraq) / Egypt (standalone from FY2025). This is MECE because every dollar of Group revenue and gross profit sits in exactly one of these three IFRS 8 reportable segments, per talabat's own disclosure structure — no market is counted twice, and none falls outside the three (GCC vs non-GCC.md; Topics/Segment Reporting.md).

– Growth-signal figures: market-comparison; TLB-019, page 14; TLB-010.)* GCC's dominance is

3.3 Competitive Analysis

SWOT — with a "so what" per quadrant

The four SWOT quadrants (internal/external × favourable/unfavourable) are MECE by construction; each is populated only with disclosed evidence, and each carries an explicit conclusion, not an orphan observation.

Competitor benchmarking

3.4 Opportunities

Untapped markets / underserved categories. Grocery/retail is the clearest untapped-opportunity category in the corpus: only ~1–3% penetrated of a USD 104–150bn+ addressable category (§3.1), and the vertical growing fastest by GMV (+47% y/y FY2025) despite carrying a structurally lower take rate/margin than Food — an explicit, disclosed (0.5)–(0.7)pp Adjusted EBITDA margin drag from the product-mix shift (Topics/Grocery and Retail.md; TLB-014, page 6). This is a genuine growth-versus-margin trade-off the plan engages with directly (consistent with OPT-001's framing) rather than presenting grocery GMV growth as equivalent in value to Food GMV growth. Non-GCC (+57% y/y GMV growth in FY2025, still only 18% of Group GMV that year, rising to 21% by Q1 2026 — TLB-019, p.14; TLB-011, p.2, two dated data points presented separately, not blended) is the second clear underserved-market opportunity by the same logic — a genuinely higher-growth segment that remains structurally under-penetrated relative to its GCC counterpart's maturity (2023 monthly orders per capita: GCC 1.28x vs. non-GCC 0.13x, a ~10x gap, TLB-026, page 121; GCC vs non-GCC.md).

Traceability (partial — full table due at Section 14)

4. Value Proposition

Capital creates value only when capability deployment changes customer behaviour before financial outcomes



Source: Value Driver Tree v2; KPI Tree v2 | Classification: Analytical causal logic, not a quantified forecast

Management implication: Release capital only when each option demonstrates its next causal link through named KPIs.

Exhibit 5. Capital creates value only when capability deployment changes customer behaviour before financial outcomes.

Answer, stated first

The value this plan creates does not come from a single new capability — it comes from applying a disciplined ranking-and-monitoring layer to five value mechanisms talabat has already proven, at Group/GCC scale, but has not yet evidenced as deliberately sequenced against its 2026 investment programme. Multi-vertical engagement (13.0x vs. 3.8x order frequency between multi-vertical and food-only customers), talabat pro (20-28% frequency uplift, 26-32% retention uplift), targeted incentives (Rewards +15%, PostPaid +14%), and AI/personalisation (embedded, not a separate budget line) are each independently evidenced value drivers — this section's contribution is showing how they trace, one by one, through Investment_Relationship_Map.md's chain to GMV, revenue, gross profit, and EBITDA, and where each one still lacks Egypt/non-GCC-specific proof.

4.1 Unique Selling Proposition (USP)

The USP is not a new customer-facing product — it is evidence-ranked, stage-gated capital allocation applied to an already-committed USD175mn programme. talabat's own disclosed rationale for the programme is qualitative ("investments rather than costs... expected to offset the margin impact over time," TLB-020, page 16); no competitor benchmark or internal capital-allocation methodology is disclosed anywhere in the corpus (Topics/Capital Allocation and Investment Governance.md). The USP is therefore comparative in kind, not degree: a transparent, evidence-traceable allocation framework where every recommended dollar amount resolves to a named Fact, Decision, or labeled Assumption — versus the undisclosed status quo. Argued narratively (per DEC-008, not as a quantified exhibit): the disclosed ~2.2:1 weighting toward Everyday App over Food-leadership is itself a real, evidenced management judgment (Topics/2026 Investment Programme.md) — this plan's initiative-level ranking (DEC-009) engages with, rather than second-guesses, that top-level weighting.

4.2 Value Creation Mechanisms — Each Quantified

Revenue lift — multi-vertical engagement. Multi-vertical customers show 13.0x order frequency vs. 3.8x for food-only customers (Topics/Multi-Verticality.md; Value_Driver_Tree_v2.md N-14), and multi-vertical GMV share has risen 68%→73%→76% across Dec'24→Dec'25→Mar'26 (Value_Driver_Tree_v2.md N-11). This is the single strongest quantified mechanism in the corpus and the direct evidence base for OPT-001 (talabat mart densification, the primary lever for driving multi-vertical adoption).

4.3 Impact on Customer Outcomes and ROI — Calculation Logic Shown

Customer outcome, stated honestly: no absolute customer lifetime value figure is disclosed anywhere in the corpus (Topics/Customer Lifetime Value.md); the mechanisms above are evidenced as relative uplifts and GMV-share trends, not an absolute CLV number this section can cite. Topics/Customer Economics.md's clearest quantified customer-outcome proxy is the >4x gap in monthly spend between multi-vertical and food-only customers — this is the number this plan uses to argue customer-level value creation, not an invented CLV figure.

Traceability (partial — full table due at Section 14)

5. AI Technology and Development

Answer, stated first

talabat does not need to build a new AI capability to execute this plan's core recommendation — the customer-facing AI capability (recommendation ranking, offer timing, dispatch optimisation) it will scale via the 2026 investment programme

is already built, licensed, and generating a disclosed, growing EBITDA contribution (>USD14mn FY2024 → >USD30mn FY2025). What this plan adds is a second, distinct AI-enabled capability that talabat does not yet disclose having: a structured decision-support layer for ranking and monitoring how the investment programme itself gets allocated. This section describes both, keeps them clearly separate, and does not invent a budget figure for either that the corpus doesn't disclose.

5.1 Description of AI Technology

talabat's embedded AI (scaled by the investment programme): machine-learning-driven personalisation (item-level and cuisine-level recommendation ranking), offer/promotion timing (triggering talabat pro and Rewards offers), and logistics dispatch optimisation — described in the corpus as producing "AI-driven logistics improvements" credited with delivery-time and cost-efficiency gains (TLB-001, page 10) and personalisation credited with timing loyalty offers "at the right time for customers" (Topics/AI.md; Topics/Recommendation Systems.md).

5.2 Build / Buy / Partner Decision — Structured Options Analysis

talabat's embedded AI capability: already a partner arrangement — talabat's technology stack is licensed from Delivery Hero SE (Entities/Technology_Platforms.md), the parent company, giving talabat access to Group-level AI capability without a standalone build. The 2026 investment programme's Everyday App opex allocation (~USD75mn) funds scaling this existing capability (more markets, more inventory, deeper personalisation), not building a new one from scratch — the build/buy/partner decision for the underlying technology was effectively made before this programme existed, and this plan does not propose revisiting it.

5.3 Proprietary Algorithms, Data Models

No proprietary algorithm or patent is disclosed anywhere in the corpus for talabat's embedded AI capability — it is described functionally (what it does: ranking, timing, dispatch), not architecturally (Topics/AI.md's Open Questions). This plan's decision-support system is not a novel algorithm either — it is a structured evidence-ranking and citation-traceability process (vault/Architecture/Decision_Management_Layer.md's 11-criterion framework; vault/Architecture/Geographic_Evidence_Rules.md's tagging discipline), not a proprietary model. Neither technology in scope here claims IP this plan can quantify or protect.

5.4 Infrastructure and Tools

talabat's own infrastructure is not itemized in the corpus beyond the Delivery Hero licensing relationship. This plan's decision-support system runs on Claude Code / Cowork (documented in MEMORY.md's LLM decision log), with the vault-based knowledge/decision/forecast architecture described throughout vault/Architecture/. MCP note: this project's own MCP integration decision remains undecided (MEMORY.md) — this section describes the OS's actual tooling (agentic subagents invoked via the pipeline in Business_Plan_Generation_Pipeline.md) honestly, without fabricating an MCP story.

5.5 Scalability and Adaptability

talabat's embedded AI capability has already demonstrated multi-market scalability (live across 6 of 8 countries for talabat pro personalisation by end-2024, excluding Egypt and Iraq — TLB-012, TLB-013, TLB-015, TLB-016, TLB-026; Value_Driver_Tree_v2.md N-08/N-13) — the investment programme's opex allocation extends this further, not builds new scalability. This plan's own decision-support system is currently scoped to five candidate Investment Options and, as of this drafting round, all 14 Business Plan sections — genuinely narrow relative to a future multi-programme or multi-year scope. Its adaptability is structural, not yet demonstrated at scale: the same Investment Option schema, 11-criterion comparison, and Geographic Evidence Rules discipline would extend to additional candidate initiatives or future investment programmes without a redesign, but this has not been tested against a larger option set.

5.6 Research and Development Plans

For talabat's embedded AI: Strategic/AI Opportunities.md names talabat's own stated Group roadmap items (evidenced intent) — this plan does not add new customer-facing AI R&D beyond what talabat has already disclosed, since the programme's opex allocation is described as scaling, not inventing, capability. For this plan's decision-support system: the two largest open R&D items are (1) resolving the cross-market allocation question (problem 3 in Section 2.2 — a disclosure gap in the corpus itself, distinct from DEC-009, which is already resolved and settles a different question: initiative-level funding sequence and allocation-range presentation within each bucket, problem 2) as better country-level evidence becomes available, and (2) instrumenting the KPI Tree's Governance family (currently 7/7 newly-instrumented KPIs with no baseline) so future allocation rounds can be evaluated against real performance data rather than constructed ranges alone.

Traceability (partial — full table due at Section 14)

Claim	Source
AI-driven logistics improvements credited with delivery-time/cost gains	TLB-001, page 10
talabat's stack licensed from Delivery Hero SE	Entities/Technology_Platforms.md
AI/personalisation EBITDA contribution USD14mn+→30mn+	Topics/AI.md; Value_Driver_Tree_v2.md N-09
N-06 not a discrete dollar line in either bucket	Value_Driver_Tree_v2.md N-06
talabat pro live in 6/8 countries by end-2024 (excl. Egypt, Iraq)	TLB-012, TLB-013, TLB-015, TLB-016, TLB-026; Value_Driver_Tree_v2.md N-08/N-13
MCP decision undecided	MEMORY.md

6. Business Model and Revenue Streams

Answer, stated first

talabat's revenue is not four independent product lines but a single GMV-times-blended-take-rate system (~38–41% of GMV across FY2024–2025) in which four disclosed fee-type components — commission fees, delivery & service fees, subscription fee & other income, and advertising & listing fees — respond to different underlying volume, penetration, or pricing drivers, and are MECE by construction of talabat's own Management Revenue reporting convention (every dollar of revenue is exactly one of the four, net of a fifth, contra-revenue "vouchers and other discounts" line). The capital-allocation-relevant finding is that the fastest-growing, most under-penetrated-relative-to-its-own-stated-benchmark lines — subscription (driven by talabat pro penetration) and advertising (driven by Partner demand for a customer base that must first be engaged) — are also the lines this plan's ranked candidate options (OPT-002 talabat pro acceleration; OPT-004 advertising monetization gap closure) are best positioned to move, while G&R's structurally lower take rate/margin (the reason OPT-001's talabat mart densification carries an explicit, disclosed EBITDA drag) means multi-vertical GMV growth is not uniformly as valuable per dollar as the headline GMV figure implies. At the per-customer level, the >4x monthly-spend gap between food-only (AED 194 / USD 53/month, 3.8 orders) and multi-vertical (AED 814 / USD 222/month, 12.8 orders) customers is the clearest disclosed unit-economics argument for the plan's cross-sell-weighted allocation. One open reporting question — the "Subscription fee & Other Income" combined line's decomposition into a true subscription component and a tMart-dominated "other income" component (ASM-013/DEC-006) — is carried forward exactly as already resolved, not re-litigated here.

6.1 Revenue Generation Methods — MECE Streams Traced to the Value Driver Tree

Why this is MECE. talabat's own "Management Revenue" reporting convention already decomposes total revenue into exactly four fee-type lines, net of one contra-revenue line — every disclosed revenue dollar belongs to exactly one of the five rows below, no overlap, no gap (Group; TLB-001, page 27; TLB-002, page 20):

Stream	FY2024	FY2025	Growth	Primary driver	Value Driver Tree node(s)
Commission fees	USD 1,062mn	USD 1,297mn	+25% → +22%	Food-vertical order volume (agent-model marketplace transaction — talabat earns commission without owning inventory); rate under some disclosed downward pressure	N-05, N-10 (Food-leadership capability), N-14/N-16 (order frequency), N-26 (Food GMV)

Stream	FY2024	FY2025	Growth	Primary driver	Value Driver Tree node(s)
Delivery & Service fees	USD 696mn	USD 859mn	+29% → +24%	Order count and basket composition; structurally eroded per-order by talabat pro's free-delivery-a-benefit even as pro raises order frequency	N-08 (pro adoption), N-14/N-16
Subscription fee & Other Income (combined line — see decomposition below)	USD 952mn	USD 1,397mn	+44% → +47%	talabat pro subscriber-penetration (GMV share 32%→49%, Q1 2025→Q1 2026) plus tMart's near-95%-take-rate owned-inventory "other income" component	N-07 (G&R product-mix shift), N-08, N-20/N-21/N-22 (retention/CLV uplift)
Advertising & listing fees	USD 246mn	USD 323mn	+27% → +32%	Partner demand for visibility — a lagging function of how engaged/frequently-ordering the customer base is, not an independent driver	N-09 (AI/personalisation), N-29, N-11/N-17 (multi-vertical usage), OPT-004
less: Vouchers and other discounts (contra-revenue)	—	USD (120)mn	—	Funded promotional /incentive spend	N-12 (targeted incentives)

6.2 Pricing Strategy

talabat's pricing is not one strategy but four distinct pricing mechanisms, one per revenue stream, each priced differently and each subject to a different constraint:

- Commission (agent-model, negotiated/rate-set) — talabat sets a commission rate on Food and Local
- Delivery & service fees (usage-based) — priced per order/basket, with talabat pro's free-delivery

6.3 Recurring vs. One-Time Revenue Streams

Reframing note, consistent with Sections 2/3's own precedent. The template's language ("subscription fees, licensing, service contracts, custom AI solutions") describes a SaaS/professional-services revenue mix that does not map cleanly onto talabat's own disclosed marketplace model — talabat discloses no "one-time" revenue line in the traditional sense (no license sale, no setup fee, no custom-project contract). This section states that mismatch explicitly rather than forcing an artificial "one-time" category the corpus does not support. The MECE distinction talabat's own disclosure actually supports is:

- Contractually recurring: subscription fees only (talabat pro membership) — the smallest of the four
- Recurring-with-usage (repeats every transaction, no contractual lock-in): commission fees, delivery

6.4 Unit Economics — What One Customer and One Order Earn and Cost

Per-customer, per-month (the clearest disclosed unit-economics differentiator in the corpus). Food-only customers spent AED 194 (≈ USD 53)/month at 3.8 orders, while multi-vertical (Food + G&R) customers spent AED 814 (≈ USD 222)/month at 12.8 orders — a >4x monthly-spend gap tied to a ~3.4x order-frequency gap (Group, September 2024 basis; TLB-026, pages 91, 122). This is the single strongest per-customer economic argument in the corpus for the plan's cross-sell-weighted allocation (OPT-001, OPT-002) over an allocation that treats every incremental order the same regardless of vertical mix — a customer who becomes multi-vertical is worth roughly four times as much per month, not a marginal increment.

6.5 The Proposed AI Decision-Support Layer's Own Revenue Treatment (cross-reference)

Consistent with Section 2.5's framing: the AI-enabled capital-allocation decision-support system this plan proposes is an internal capability, not a separately-sold product, and therefore does not add a new revenue line to the MECE breakdown above. Its value is realized entirely through the four existing streams' better allocation — specifically, by helping direct capital toward the mechanisms (OPT-002 talabat pro; OPT-004 advertising) shown above to sit closest to the under-penetrated, fastest-growing lines, rather than by monetizing the decision-support layer itself. This is a deliberate scope choice carried forward from Section 2, not a gap in this section's own analysis.

Traceability (partial — full table due at Section 14)

7. Marketing and Sales Strategy

Answer, stated first

talabat's own disclosed marketing posture is not price-led acquisition but customer-value-proposition (CVP) investment — management states explicitly that Food-leadership capital goes toward CVP "rather than matching competitor discounts and vouchers" (Group; TLB-020, page 16) — and this plan's marketing and sales strategy for the 2026 investment programme adopts that same positioning by design, not by default: acquisition and retention spend should be concentrated on the evidenced, product-embedded mechanisms already ranked in Strategic/Customer Retention Drivers.md (talabat pro, multi-vertical cross-sell, Rewards, PostPaid, Family Plan) rather than broad vouchering, and should be sequenced across markets the way talabat itself already sequenced FinTech (GCC-first, Egypt as the first disclosed non-GCC follow-on) rather than assumed uniform across all eight markets at once. The genuine, unresolved gap this section must handle openly is that no funnel-stage conversion data (awareness → trial → subscription) exists anywhere in the corpus at any geography level — the customer-journey sequence (Food → Grocery & Retail → talabat pro sign-up) is real, disclosed evidence, but the percentage of customers who move between those stages is not disclosed anywhere; this section states that gap plainly and uses explicitly-labeled synthetic figures, never corpus-cited ones, to illustrate funnel economics as a framework. Partnerships and retention tactics both draw on the same ranked mechanism set: talabat pro, Rewards, PostPaid, and the co-branded bank cards, positioned as product-embedded loyalty investment, with OPT-002 (talabat pro acceleration in Egypt/Iraq) and OPT-003 (Food-leadership CVP targeted at UAE/ Kuwait/Qatar) as the plan's own ranked, evidence-grounded candidates for where marketing/sales-adjacent capital should go first.

7.1 Positioning — What talabat Is, and Deliberately Is Not, Competing On

The stated positioning. talabat frames itself as transitioning "from a multi-vertical food-delivery platform to the region's Everyday App" (Group; TLB-020, page 16, cited in Section_02_Business_Description.md §2.1), a narrative arc management itself traces across three successive brand statements — "Delivering food to your doorstep" (2016) → "Delivering food and a lot more" (2021) → "It's not just talabat app, it's your app" (2025) (Group; TLB-016, page 7; Topics/Customer Journey.md). Against a competitive backdrop management itself names as intensifying — "a more competitive environment (new entrants and incumbents)" is a stated FY2026 headwind (Group; TLB-019, page 18; Topics/Competition.md), with non-high-value customer M1 retention having declined 4% year-over-year specifically in UAE, Kuwait, and Qatar (market-comparison: UAE, Kuwait, Qatar; TLB-019, page 9) — talabat's own stated response is not to match on price.

7.2 Customer Acquisition — Channels, Funnel, and Journey Mapping

Disclosed channels (Fact, Group level). The corpus discloses acquisition and retention spend by cost category, not by discrete marketing channel (no disclosed breakdown of paid search, social, in-app referral, or offline spend exists anywhere in the corpus). What is disclosed: total Customer Acquisition and Retention Costs (CARC) rose from USD 89mn (1.5% of GMV, 2023) to USD 155mn (1.6% of GMV, 2025), of which USD 89mn (2025) was talabat-funded vouchering specifically (Group; TLB-001, page 28; TLB-002, page 21; Topics/Promotions.md); separately, Partner-funded savings (not a talabat marketing-spend line, but a channel-adjacent customer-facing incentive) reached an all-time-high 7% of GMV in Q1 2026, totaling more than AED 1,567.7mn over a trailing 12-month period (Group; TLB-020, page 7; TLB-023, page 6). So what for the P&L: because talabat's own cost taxonomy already separates "vouchering" (a discount cost) from the CARC total, this plan's acquisition strategy inherits that same discipline — proposing that any 2026-programme acquisition spend be tracked and evaluated the same way, as an investment with a measurable frequency/retention return, not merely a discount cost (Topics/Promotions.md Business Implications).

Funnel stage	Anchor to real, disclosed evidence	Illustrative conversion (stage-to-stage)
Awareness → App download / first visit	Not disclosed (no stated basis)	[SYNTHETIC — illustrative estimate, not derived from talabat's disclosed data]
Download → Trial (first order placed)	Not disclosed (no stated basis)	[SYNTHETIC, illustrative: ~25–35%]

Funnel stage	Anchor to real, disclosed evidence	Illustrative conversion (stage-to-stage)
Trial → Reaches ~6-order personalisation threshold	Real, disclosed mechanical threshold (Group; TLB-002, p.15) — but no disclosed % of trialists who reach it	[SYNTHETIC, illustrative: ~15–25% of first-time orderers]
6-order threshold → Multi-vertical adoption (Food + G&R)	Real, disclosed journey sequence (Group; TLB-013, p.6) — but no disclosed conversion %	[SYNTHETIC, illustrative: ~30–40%]
Multi-vertical → talabat pro sign-up	Real, disclosed journey terminus (Group; TLB-013, p.6) — but no disclosed conversion %	[SYNTHETIC, illustrative: ~10–20%]

7.3 Partnerships

Restaurant/retail Partners — the disclosed core of the partnership base. Food-leadership's stated partner-side investment is explicit: "on the partner side, we invest in retaining, winning back, and acquiring high-demand food partners... reflected in commission rate investments" (Group; TLB-020, page 16; TLB-014, page 19; Topics/Food Leadership.md) — a qualitative mechanism with no disclosed partner-retention percentage or count anywhere in the corpus (vault/Forecasts/Value_Driver_Tree_v2.md N-10). This is directly relevant to marketing/sales strategy because Partner economics fund a meaningful share of customer-facing promotional value: Partner-funded savings reached an all-time-high 7% of GMV in Q1 2026 (Group; TLB-020, page 7) — distinct from, and larger in scale than, talabat's own voucher spend.

7.4 Retention — Support, Community, and Product Enhancement

The ranked mechanism set, restated for the sales-and-marketing lens. Strategic/Customer Retention Drivers.md's own synthesis ranks five product-embedded retention levers by strength of corpus evidence: (1) talabat pro — a 20-28% order-frequency uplift and 26-32% retention uplift versus matched non-subscribers, and a 136% gross-profit-per-customer uplift in the 30 days following subscription (GCC + Jordan, six markets live before December 2024; explicitly excludes Egypt and Iraq by name; TLB-001, page 18; TLB-015, pages 78, 104; TLB-019, page 11); (2) multi-vertical engagement, which amplifies pro's effect rather than operating independently — mono-vertical subscribers show a +16pp M1 retention lift versus mono-vertical non-subscribers, multi-vertical subscribers show +20pp (GCC + Jordan; TLB-019, page 10); (3) talabat Rewards, an 18%-adoption, >15%-order-frequency-uplift-within-30-days mechanism that talabat's own language links directly to "reducing churn" (Group; TLB-026, page 134; TLB-015, page 79); (4) talabat PostPaid, a 14% order-frequency-increase-post-adoption mechanism described as "accretive to order frequency and customer retention" (Group; TLB-001, pages 9, 18); (5) Family Plan, a >60% retention premium over solo-plan subscribers, on the narrowest evidence base of the five (single document, single household-size segment, no adoption-rate or country breakout) (Group; TLB-018, page 6).

7.5 Where This Positions the Plan's Own Ranked Options

Restated plainly: this section's positioning (CVP over discounting), acquisition approach (product-embedded journey stages plus a labeled synthetic funnel model, sequenced on the disclosed GCC-first-then- Egypt FinTech precedent), partnership base (restaurant/retail Partners, banking co-brand cards), and retention tactic ranking (pro > multi-vertical > Rewards/PostPaid > Family Plan, AI as infrastructure) are not new claims invented for this section — they are the marketing-and-sales-strategy expression of the same evidence base DEC-009 already used to rank OPT-001–OPT-005. Specifically: OPT-002 (pro acceleration, Egypt/Iraq) and OPT-003 (Food-leadership CVP, UAE/Kuwait/Qatar) are this section's two most directly relevant candidates — DEC-009's resolved funding-sequence recommendation (not merely its earlier, purely descriptive 11-criterion tiering) names both jointly as the near-term priority pair, precisely because their profiles are complementary: OPT-002 is the higher-uncertainty, highest-reversibility candidate this section's own Egypt-first-then-Iraq sequencing (§7.2) is built around, best suited to a bounded pilot; OPT-003, by contrast, is an already-evidenced, already-operating mechanism that DEC-009 and OPT-003's own option record recommend funding now, as a targeted budget weighting within the existing Food-leadership bucket — explicitly "not a pilot." Section 12's Three Horizons framing places both in Horizon 1 on this basis (one piloted, one funded directly), not because every near-term option must be piloted before scaling.

Traceability (partial — full table due at Section 14)

8. Operations Plan

Answer, stated first

talabat does not need to build new operational infrastructure to execute this plan's recommendation — its technology and delivery infrastructure is already licensed and operating at Group scale through a disclosed set of inter-company services agreements with parent Delivery Hero SE (the Global Licensing and Services Agreement and its category-specific variants), and its logistics/dispatch stack already runs across all eight markets. What this plan's own 7S alignment check finds is a genuine, honestly-named organizational gap, not an infrastructure one: talabat's disclosed operating model is strong on Strategy (the Everyday App/Food-leadership split and the CVP-over-discounting positioning are clear and consistently communicated) and Style (a recognisable, repeated investor-communication cadence), adequate but concentrated on Structure and Staff (governance and named executive roles are visible but centred on Delivery Hero SE-overlapping seats, with two unexplained recent leadership transitions), and not ready on Systems — no internal capital-allocation committee, approval threshold, or stage-gate process for this specific USD175mn programme is disclosed anywhere in the 29-document corpus (Topics/Capital Allocation and Investment Governance.md) — and shows a real Skills gap, since no named data-science, capital-allocation, or portfolio-management function appears in the executive roster this corpus discloses. This section names those misalignments honestly and proposes closing the Systems gap with a stage-gate mechanism this OS explicitly labels as its own recommendation, gated for human approval, not a description of how talabat actually runs its own capital-allocation process today.

8.1 Infrastructure Needs

The technology layer is inherited, not built. talabat's underlying infrastructure runs on capability licensed from Delivery Hero SE rather than owned or engineered standalone by talabat itself, delivered through a disclosed Global Licensing and Services Agreement (GLSA): members of the Group in Bahrain, Egypt, Iraq, Oman, Jordan, and Qatar receive a defined "Central Value Basket" of data-management and communication tools, logistics tools (including rider-recruitment tools), customer-management tools, audit tools, quick-commerce tools (catalogue and assortment intelligence, purchase-management tools, supplier portal, inventory-management tools, store management, rider-management tools, advertising and promotion tools), FinTech tools (wallets, payment and fraud-protection solutions), and business-management services spanning sales/operations and international marketing, including "customer loyalty and subscription" support (market-comparison: Bahrain, Egypt, Iraq, Oman, Jordan, Qatar; TLB-026, pages 152-153) — with parallel, separately-named GLSA agreements covering UAE and Kuwait specifically (implied by the corpus's reference to "the GLSA, the GLSA Kuwait and the GLSA UAE" as distinct instruments; TLB-026, page 156). talabat pays Delivery Hero arm's-length compensation for this basket, calculated per specific service line (TLB-026, page 153). So what: this is the concrete evidence base for Section 5.2's build/buy/partner finding — the "partner" arrangement is not a general description, it is a named, priced, multi-market services contract that already covers most of what OPT-001 through OPT-005 would draw on operationally (inventory-management tools for OPT-001; loyalty/subscription support for OPT-002; advertising/promotion tools for OPT-004; the shared data layer for OPT-005).

8.2 Development and Maintenance Workflow

What is disclosed: a contractual services structure, not an internal software development lifecycle. No document in the corpus describes talabat's internal engineering process (sprint cadence, release process, environment/staging structure, or code-ownership model) — this is a genuine gap, stated openly rather than invented. What is disclosed is how the underlying capability is provided and paid for: the GLSA (§8.1) sets the terms on which Delivery Hero supplies and talabat "localises and uses" the Central Value Baskets (Group/market-comparison; TLB-026, pages 152-153), the CQCA set equivalent terms for tMart-specific tools through 31 December 2024, since when the same tools are delivered via GLSA sub-licensing instead (Group; TLB-026, page 156), and a further Kitchens Services Agreement (effective 1 January 2022, amended 29 October 2024) covers the Kitchens business specifically in the UAE, Bahrain, Kuwait, Qatar, and Jordan (market-comparison: UAE, Bahrain, Kuwait, Qatar, Jordan; TLB-026, page 156). All three are governed by German law, with disputes resolved in German courts (GLSA: TLB-026, page 153; CQCA and Kitchens Services Agreement: TLB-026, page 156) — a detail worth naming because it confirms these are formal, arm's-length inter-company contracts, not an informal shared-services arrangement. So what: "maintenance" of talabat's core technology, for the purposes of this programme, is contractual and inter-company, not a talabat-internal release cycle this plan can describe or optimise directly — any recommendation this plan makes about accelerating a technology-dependent option (e.g. OPT-001's inventory tooling, OPT-005's model tuning) implicitly depends on Delivery Hero's own service delivery under these agreements, a dependency this plan names rather than assumes away.

8.3 Key Team Roles and the Skills Gap vs. Today

Named roles the corpus discloses that this programme would plausibly draw on. talabat's executive roster includes several roles directly relevant to executing and governing this programme: Khaled Alfakesh, Chief Financial Officer since 2016, the most visible existing sponsor for a capital-allocation discipline function; Pedram Assadi, Chief Operations Officer; Wassim Makarem, SVP Grocery & Retail, described as "at the forefront of driving talabat's regional quick commerce initiatives" — the natural existing owner for OPT-001; Stefano Vecchio, VP People & Strategy, who also oversees "New Ventures" including talabat's loyalty programmes — the natural existing owner for OPT-002; and a

governance- adjacent cluster — Mohamd Abu Amara (Head of Governance, Risk & Compliance), Oussama El Kadri (Head of Internal Audit), and Abdullah AIGHrawi (VP Legal, GRC & Board Secretary, the DFM buyback notice signatory) (Group; Entities/Executives.md, TLB-001, TLB-021).

8.4 7S Alignment Check

The corpus does not contain a pre-built 7S analysis for talabat — this is genuine synthesis, built here for the first time against the evidence assembled above and in the Topic Notes cited throughout this plan. Each "S" is rated Ready, Partially ready, or Not ready, with the specific evidence and gap-closing action named — no unsupported rating.

S	Rating	Evidence	Gap-closing action
Strategy	Ready	The Everyday App (~USD120mn) / Food-leadership (~USD55mn) split and the CVP-over-discounting positioning are clearly disclosed and consistently communicated (Section_02_Business_Description.md §2.1; Section_07_Marketing_and_Sales_Strategy.md §7.1; TLB-020, page 16)	None required — this is the plan's strongest evidenced "S"
Structure	Partially ready	Board dominated by Delivery Hero SE officers (Vandepitte/Chair also DH SE COO; Krause also DH SE General Counsel; Popp also DH SE interim CFO — Corporate Structure.md, TLB-026); Egypt operates through two dedicated legal entities (Delivery Hero Egypt SAE, Delivery Hero Dmart Egypt LLC) with no disclosed local board or decision rights (Corporate Structure.md Open Questions)	This plan's recommendation is presented for Group-level approval, consistent with the visible governance concentration, rather than assuming a country-level approval path exists (Decision-Making Process.md Strategic Implications)
Systems	Not ready — the largest gap	No internal capital-allocation committee, approval threshold, or stage-gate process for the USD175mn programme is disclosed anywhere in the corpus (Capital Allocation and Investment Governance.md)	This OS's own proposed mechanism, not a description of talabat's actual process: a lightweight stage-gate — each OPT- option reviewed against its own named stage gate (already specified per option, e.g. OPT-002's 2-quarter Egypt pilot checkpoint, OPT-001's margin-drag checkpoint) before scale-up funding, tracked against KPI_Tree_v2.md's Governance family (G1-G7, all seven newly-instrumented, none with a corpus baseline) — gated for human/team approval before any capital moves, per the Responsible-AI principle Section 11 develops further
Shared Values	Partially evidenced	Management's own stated framing — "these are investments rather than costs" (TLB-020, page 16) and CVP-over-discounting (TLB-020, page 16) — is a real, quotable value statement; no equivalent explicit value statement about AI ethics or data governance is disclosed beyond the DTA's compliance mechanics (§8.5)	This plan's Section 11 (Responsible AI) is the designated place to state the value this OS itself operates on (human-approval-gated recommendations, never a false-precision automatic decision) explicitly, since talabat's own disclosure does not state an equivalent value for its internal allocation process

S	Rating	Evidence	Gap-closing action
Skills	Not ready	No named data-science, capital-allocation, or portfolio-management function in the disclosed executive roster (§8.3); OPT-004/OPT-005 both name a specific, currently-unconfirmed capability gap (Partner-facing ad sales; per-market AI model tuning)	Locate the proposed Systems function (above) near the existing CFO/GRC structure rather than inventing a new seat; pilot OPT-005's per-market model-tuning gap in Egypt specifically, per OPT-005's own pilot recommendation
Style	Ready	A recognisable, repeated communication cadence — Capital Markets Day for medium-term targets, Annual Report for the year's strategic objective, quarterly results for dollar figures (Decision-Making Process.md "venue pattern") — and consistent language ("investments not costs," CVP-over-discounting) across TLB-014/019/020	None required
Staff	Partially ready	Governance-adjacent named roles exist (GRC, Internal Audit, Legal/Board Secretary — §8.3) — but two leadership transitions are recorded in the corpus without any narrated explanation: CEO Tomaso Rodriguez → Toon Gyssels (between TLB-018, Aug 2025, and TLB-009, Feb 2026) and board seat Muhammad Hussain Ghati Al Jbori → Abdul Wahab Al-Halabi (between TLB-004/TLB-026 and TLB-008) (Corporate Structure.md, Entities/Executives.md)	Named explicitly here as a continuity caveat and carried forward into Section 10 (Risk Analysis) as an organizational risk, not resolved or explained by this plan, since no source in the corpus does so either

8.5 Security and Compliance Measures

A disclosed security incident, not a hypothetical risk. In December 2022, talabat was hacked by an external attacker based in Norway, who gained access to the personal data of 144,469 customers in one of talabat's markets (market unnamed in the disclosure); talabat informed the competent data-protection regulator, which opened an investigation, and talabat paid a USD 150,000 penalty (Group, specific market not disclosed; TLB-026, page 49). The same risk-factor disclosure states plainly that because talabat's technology infrastructure — including payment solutions and data storage — depends on Delivery Hero's systems, talabat is "vulnerable to any security breaches or data protection issues that may occur at the parent company level" (Group; TLB-026, page 49) — a direct, disclosed consequence of the same partner/licensed-infrastructure model described in §8.1.

Traceability (partial — full table due at Section 14)

9. Financial Plan

Answer, stated first

talabat's own FY2026 guidance already prices in the 2026 investment programme's cost: GMV growth of 11-14% cFX, Revenue growth 14-17% cFX, Adjusted EBITDA of 4.4-4.8% of GMV (down from 6.0% in Q4 2025), and Free Cash Flow of 3.2-3.6% of GMV. This plan's financial case does not construct a separate, more precise projection than talabat's own guidance — per DEC-008, it presents that guidance as the headline base case, with upside and downside scenarios bounding it, and explicitly does not attempt to isolate a return figure for either investment bucket, because no such figure is disclosed anywhere in the corpus.

9.1 Value Driver Tree — Linking the Investment Programme to Financial Outcomes

The governing chain (vault/Knowledge/Investment_Relationship_Map.md; formalized as Value_Driver_Tree_v2.md's 45 nodes): Investment → capability deployment → adoption/operational change → customer/partner behaviour → order frequency → multi-vertical usage → basket/AOV → retention/CLV → GMV → revenue → gross profit → EBITDA → cash flow. Every node from "customer/partner behaviour" onward that connects the programme to a financial outcome is an Assumption node, not a Fact node — the corpus discloses the programme's cost (the EBITDA margin bridge) but not a quantified return at any point in this chain (Value_Driver_Tree_v2.md, "Central evidence-gap this tree does not paper over").

9.2 Initial Investment and Operational Costs

Disclosed, not decided by this plan: USD175mn total, Board-approved February 2026, fully funded by internal cash — ~USD120mn Everyday App (~USD75mn opex + ~USD45mn capex) + ~USD55mn Food-leadership (vault/Decisions/Investment_Portfolio_Register.md). This plan's own analytical sub-allocation (OPT-001–005, ranged in the Portfolio Register, ranked and sequenced in DEC-009, status: approved 2026-07-23 — middle path) is referenced here narratively, not shown as a headline financial exhibit, per DEC-009's own mandatory placement rule (the constructed dollar ranges belong in a Section 14 appendix/sensitivity exhibit, each carrying the sentence "This is an Agentic OS analytical recommendation, not a disclosed talabat allocation," never here): the five options span both buckets plus two cross-cutting enabling-capability candidates, are explicitly non-additive against the USD175mn total, and do not claim to allocate 100% of either bucket (Investment_Portfolio_Register.md's stated caveats). Recommended funding sequence (per DEC-009): near-term priority to OPT-002 (talabat pro Egypt/Iraq — cheap, fast, reversible) and OPT-003 (Food-leadership GCC-3 — tightly evidenced); OPT-001 (dark-store densification) and OPT-005 (AI/personalisation) continue at pace; OPT-004 (advertising gap) is a longer-horizon build.

9.3 Revenue Forecasts — Base, Upside, Downside (FY2026)

Scenario	GMV growth (cFX)	Revenue growth (cFX)	Adj. EBITDA margin	FCF margin	Basis
Base	11-14%	14-17%	4.4-4.8% of GMV	3.2-3.6% of GMV	talabat's own disclosed FY2026 guidance (TLB-020, TLB-014)
Upside	Toward/above 14%	Toward/above 17%	Toward 4.8%+	Toward/above 3.6%	Everyday App/Food-leadership mechanisms (multi-vertical, pro, advertising) compound faster than guidance embeds — Scenarios_v2.md upside case
Downside	Toward/below 11%	Toward/below 14%	Toward/below 4.4% (Q1 2026 actual already 4.8%, -9% y/y in absolute EBITDA dollars)	Toward/below 3.2%	Margin compression deepens, behaviour-change trends moderate faster than guidance assumes — Scenarios_v2.md downside case

All three rows are Group-level, per DEC-008. No bucket-level, market-level, or initiative-level breakout of these figures is shown here — none is disclosed, and constructing one would be exactly the false-precision error DEC-008 was escalated to prevent.

9.4 Break-Even Analysis and P&L Snapshot — What Can and Cannot Be Shown

A conventional break-even/payback analysis requires a return figure to compare against the USD175mn outlay. No such figure exists in the corpus for either bucket or any candidate initiative — confirmed independently by Investment_Relationship_Map.md, EBITDA.md, and Capital Allocation and Investment Governance.md. What can be shown honestly: the disclosed cost side, as a P&L snapshot — Adjusted EBITDA margin moving from 6.0% (Q4 2025 actual) to a realized 4.8% (Q1 2026 actual) to a 4.4-4.8% guided range (FY2026), a temporary, deliberate step-down management itself frames as investment, not deterioration (TLB-020, page 16). This plan does not construct a break-even date or payback period this evidence cannot support.

9.5 Key Assumptions Register (the numbers this case depends on)

All fourteen listed are status: Approved in vault/Decisions/Assumptions_Register.md (resolved 2026-07-23, alongside DEC-009) — citable as settled inputs per the register's own rule. Listed here for transparency, exactly as the pipeline requires:

9.6 Funding Requirements, Allocation, and Monetization Strategy

Funding requirement: none beyond what talabat has already committed — the USD175mn programme is fully funded from internal cash (TLB-020, page 16); this plan does not request incremental capital beyond the disclosed envelope. Allocation: governed by the disclosed ~120mn:55mn bucket split (fixed) and, within it, the five candidate options this plan has ranked and sequenced per DEC-009 (approved 2026-07-23, middle path): OPT-002 and OPT-003 as near-term priority, OPT-001 and OPT-005 continuing at pace, OPT-004 as the longer-horizon build — stated narratively as the plan's actual recommendation; the underlying ASM-029–033 dollar ranges are never a headline figure here, and appear only in Section 14's labeled sensitivity exhibit. Monetization: not applicable in the conventional sense — this is a capital-allocation decision-support exercise, not a revenue-generating product; its "monetization" is the incremental GMV/EBITDA the better-allocated USD175mn is expected to generate relative to an unranked baseline, a claim this section deliberately does not quantify, per §9.1's evidence-gap finding.

Traceability (partial — full table due at Section 14)

10. Risk Analysis




Exhibit 7. Three high-priority risks require management attention before capital release.

Answer, stated first

This plan's largest risk is not algorithmic — it is that talabat is funding a margin-dilutive, evidence-thin capital-allocation programme with almost no disclosed internal mechanism to correct course, layered on top of already-realized competitive, regulatory, and leadership volatility. Of the 17 material risks identified below across five MECE categories (Technical, Market, Financial, Organizational & Governance, Regulatory), three land in the probability-impact matrix's highest-severity quadrant (High probability × High impact): the absence of any disclosed internal stage-gate or approval mechanism for the USD175mn programme (ORG-01), the already-observed FY2026 Adjusted EBITDA margin step-down (FIN-01), and competitive pressure eroding Food's non-high-value customer base faster than its ~USD55mn defense budget (MKT-01). A fourth category-defining risk — the possibility that this OS's own constructed allocation ranges get mistaken for a disclosed talabat commitment (ORG-02) — is unique to this plan's own recommendation process, not a risk talabat's disclosures describe, and is named here explicitly because DEC-008 and DEC-009 were each escalated specifically to guard against it. Every risk below carries a named owner and a mitigation grounded in already-disclosed evidence or this plan's own prior pipeline output (Section 8's proposed stage-gate mechanism, the Governance KPI family, DEC-008/ DEC-009's disclosure discipline) — each is evidence-specific. This risk list is the direct evidentiary anchor Section 11 (CSR & Responsible AI) must operationalize, particularly ORG-01, ORG-02, and ORG-03.

10.1 Technical Risks

This breakdown is MECE by proximate cause, not by every downstream effect: each risk is assigned to exactly one category based on the mechanism through which it first materializes, with secondary effects in other categories cross-referenced inline rather than double-counted in the matrix (§10.6) or the risk total.

10.2 Market Risks

MKT-01 — Competitive pressure is already eroding Food's non-high-value customer retention faster than its capital allocation defends against it. The disclosed non-high-value M1 retention decline is -4% y/y, scoped specifically to UAE, Kuwait, and Qatar (market-comparison: UAE, Kuwait, Qatar; TLB-019, p.9) — Food-leadership's ~USD55mn CVP/partner-retention investment is funded to counter exactly this pressure (Topics/Food Leadership.md), yet Food still represents the majority of Group GMV (USD6.65bn vs. G&R's USD2.77bn, FY2025; Group; TLB-002, p.18) against under a third of the 2026 programme's total capital — a genuine, quantifiable tension Food Leadership.md's own Business Implications names rather than resolves. Whether Egypt or other non-GCC markets face comparable pressure is not evidenced anywhere in the corpus (Food Leadership.md Open Questions) — extending this risk beyond UAE/Kuwait/Qatar is a labeled inference, not a disclosed fact. Geography: market-comparison (UAE, Kuwait, Qatar) for the retention-decline evidence; Group for the GMV/funding comparison; inferred-applicability if extended to Egypt/non-GCC.

10.3 Financial Risks

FIN-01 — The disclosed FY2026 Adjusted EBITDA margin step-down is real, evidenced, and already in motion, with a documented trail — not a hypothetical risk. Group Adjusted EBITDA margin fell from 6.7% (FY2024) to 6.5% (FY2025); FY2026 guidance steps this down further to 4.4-4.8% of GMV, a deliberate choice funding the USD175mn programme, not a surprise (Group; TLB-020; Strategic/Strategic Risks.md). Q1 2026 actual Adjusted EBITDA margin was already 4.8%, at the guided range's low end, with Adjusted EBITDA down 9% year-on-year and Net Income down 18% year-on-year even as GMV and Revenue grew (Group; TLB-020, pp.4, 10-11). Problem_Charter.md separately frames the same Q1 2026 4.8% actual figure against the prior-year quarter specifically ("compressed to 4.8% from 6.3% a year earlier," Group; Problem_Charter.md) — a different, non-contradictory comparator (year-on-year vs. this section's quarter-on-quarter framing above) for the same underlying print. Scenarios_v2.md's downside case extrapolates this trend deepening beyond guidance's low end if the G&R margin drag (FIN-04 below) or the Everyday App/Food-leadership margin bridge's own unreconciled quarterly-vs-annual attribution prove larger than currently guided. Geography: Group.

10.4 Organizational & Governance Risks

ORG-01 — No internal capital-allocation committee, approval threshold, or stage-gate process is disclosed anywhere in the corpus. This is the single largest, load-bearing governance gap this plan works around: Topics/Capital Allocation and Investment Governance.md states plainly that no document describes how a specific initiative (country-level or category-level) moves from proposal to funded line item, or how — or whether — the Board could reallocate the USD175mn programme intra-year if early results diverge from plan. Section_08_Operations_Plan.md's 7S alignment check independently confirms Systems as talabat's least organizationally-ready dimension for exactly this reason. Without a disclosed correction mechanism, capital could keep flowing to an underperforming option (e.g., OPT-001 if its margin-drag trajectory worsens, per FIN-04) with no committed trigger to stop it. Geography: Group.

10.5 Regulatory Risks

REG-01 — Antitrust/competition-law exposure tied directly to talabat's market dominance. Ongoing Kuwait Competition Protection Authority investigations/litigation over Partner dealings, UAE Competition and Consumer Protection Department notices over Partner-subscription contract clauses, and an Iraq competition-authority inquiry (no formal complaint) are each disclosed (country-specific: Kuwait, UAE, Iraq; TLB-001, pp.32,34; TLB-002, p.26; TLB-026, pp.46-47), per Strategic/Competitive Weaknesses.md. The corpus's own language — "antitrust scrutiny... related to talabat's strong market position" (TLB-001, p.34) — makes the causal link explicit: dominance itself invited this exposure. Geography: country-specific (Kuwait; UAE; Iraq — three separately disclosed items, not a single multi-country claim).

10.6 Probability-Impact Matrix

New synthesis, built fresh against the evidence assembled above — not a pre-existing corpus artifact. Probability and Impact are each rated High/Medium/Low based on: Probability — whether the risk is already observed/realized in the corpus (High), a plausible extension of an observed trend or named but unresolved exposure (Medium), or a low-likelihood/narrow-scope item (Low); Impact — the risk's potential effect on the USD175mn programme's ability to deliver its stated exit criterion (TLB-020's "more loyal customer base") at Group scale (High), at a bucket/option level (Medium), or narrowly scoped to one option/market (Low).

10.7 Pre-Mortem: "It is FY2028, and the 2026 investment programme failed to deliver a more loyal

customer base — why?"

— **The governance gap never closed, and the plan's own numbers were mistaken for commitments

– **Margin compression deepened past the point of return, and no one could tell which line was responsible

10.8 Mitigation Strategy and Owners (per material risk)

Owners are drawn from talabat's disclosed executive roster (Entities/Executives.md) where a function plausibly exists; where no such function is named — a gap this plan states explicitly rather than inventing a role — the nearest disclosed function is named as interim owner, consistent with Section_08_Operations_Plan.md §8.3's own treatment of the same gap.

Traceability (partial — full table due at Section 14)

11. Corporate Social Responsibility & Responsible AI

Answer, stated first

This plan treats Responsible AI as risk management, not decoration, because it has to: this OS's own recommendation process creates a real risk (ORG-02, the false-precision risk DEC-008 and DEC-009 were each escalated to guard against), sits inside an organization with no disclosed correction mechanism for its own capital-allocation programme (ORG-01), and depends on leadership continuity that has already shown two unexplained transitions (ORG-03). Every commitment in this section is a concrete, already- operating control tied by name to one of these three Section 10 risks, not a generic AI-ethics statement: (1) this plan's constructed numbers are never shown as if internally verified — a control already exercised twice in practice, not just asserted, at DEC-008 (Section 9's headline case restricted to talabat's own disclosed base-case guidance) and DEC-009 (the five investment options' cost ranges confined to a labeled sensitivity exhibit, never a headline commitment); (2) every geography-crossing inference is explicitly labeled, per Geographic_Evidence_Rules.md, preventing the specific bias failure mode (misapplying GCC- measured evidence to Egypt as if disclosed there) this pivot's own root-cause analysis identified; and (3) no capital-allocation recommendation this plan makes is presented as an automatic management decision — Investment_Portfolio_Register.md's Owner section and Section_08_Operations_Plan.md §8.4's proposed stage-gate mechanism both require human/team approval before any capital actually moves, directly operationalizing Problem_Charter.md's own Responsible-AI principle. Sustainability initiatives (§11.2) are real and disclosed — EV delivery pilots in Egypt and the UAE, Group-level emissions reporting, rider training and compensation, and named community programmes — but this section is honest that no disclosed link exists between them and the USD175mn programme this plan is about, and states that gap rather than implying one.

11.1 Ethical AI Commitments Tied to Concrete Controls

This breakdown is MECE by control type, not by ethical principle in the abstract: each commitment names exactly one operating control, the Section 10 risk (or, for Control 4, the stakeholder expectation) it answers, and the evidence that the control already works in practice, not just in policy language.

Commitment	Concrete control	Section 10 risk / stakeholder expectation	Evidence the control is real, not aspirational
No false-precision capital numbers	Headline-exhibit restriction + mandatory disclosure sentence on every constructed range	ORG-02	DEC-008 (Option 1, approved); DEC-009 (middle path, approved)
Geography-crossing inferences always labeled	Nine-tag geography system; checked at citation audit (13.9) and Geographic Evidence Gate (Stage 16)	Fairness/evidentiary-integrity risk this pivot's own root-cause analysis identified	Geographic_Evidence_Rules.md; the superseded Value_Driver_Tree.md case study
No customer-level data claim	Aggregate-only inputs (Group/GCC/segment)	Stakeholder expectation (Board/investor trust, post-TECH-02/REG-04)	Problem_Charter.md; Section_08_Operations_Plan.md §8.5
No assumed-uniform AI effectiveness across markets	OPT-005 per-market model-tuning pilot, Egypt first	TECH-01	Section_08_Operations_Plan.md §8.4; Section_10_Risk_Analysis.md §10.8

11.2 Sustainability Initiatives

talabat discloses real, if partial, sustainability activity — this plan reports it honestly, including what it does not know. Four categories of disclosed activity exist in the primary corpus:

- Electric-vehicle delivery fleet. talabat operates 240 electric vehicles in the UAE (as of January
- Group-level emissions reporting. talabat reports total 2025 emissions of 875,157 tCO₂e (Scope 1:

11.3 Commitment to Responsible AI Practices and Governance Structures

The governing principle, stated at charter level and never diluted downstream. Problem_Charter.md states: "The AI does not claim to possess talabat's internal, customer-level data, and its output is not a substitute for management decision-making... every allocation recommendation in the resulting plan is a range or a staged-funding proposal grounded in disclosed evidence and explicitly labeled assumptions, gated for human approval before any capital actually moves, never a false-precision single number presented as if it were internally verified."

Investment_Portfolio_Register.md's Owner section restates this operationally: "This register is owned by the Decision Steward Agent, with capital-allocation recommendations ultimately requiring human/team approval before being presented as a Business Plan recommendation." This section is where that principle becomes a named governance structure this plan commits to, tied explicitly to ORG-01 and ORG-03.


11.4 What this section does not claim

Consistent with the discipline the rest of this plan applies, this section states three things it explicitly does not do: (1) it does not claim talabat has adopted, or has committed to adopt, any of the governance structures proposed above — they are this OS's own recommendation, pending human/team decision, exactly as Section_08_Operations_Plan.md §8.4 and Capital Allocation and Investment Governance.md state; (2) it does not claim the sustainability initiatives in §11.2 are funded by, or governed under, the USD175mn 2026 investment programme — no such linkage is disclosed; (3) it does not claim to have resolved ORG-01 or ORG-03 — both remain open risks in Section_10_Risk_Analysis.md's probability-impact matrix even after this section's proposed mitigations, because a proposed mechanism, not yet adopted, does not retire a risk that depends on adoption.

Traceability (partial — full table due at Section 14)

12. Implementation Plan (Three Horizons)

Three proof gates sequence quick wins before broader scale commitments



Source: Section 12, DEC-009, Investment Options Register | Classification: Analytical implementation horizon; no new budget
Management implication: H2 and H3 capital release depends on evidence generated in the prior horizon.

Exhibit 6. Three proof gates sequence quick wins before broader scale commitments.

Answer, stated first

The five candidate Investment Options this plan has identified do not all belong at the same starting point. DEC-009's underlying 11-criterion comparison groups OPT-001 (talabat mart densification) and OPT-005 (AI/personalisation scaling) as the strongest-evidence, broadest-applicability pair — continue-at-pace items requiring no new pilot. But DEC-009's resolved funding-sequence recommendation is not a simple restatement of that evidence tiering: it names OPT-002 (talabat pro acceleration, Egypt/Iraq) and OPT-003 (Food-leadership CVP targeting, UAE/Kuwait/Qatar) as the near-term priority pair, precisely because their profiles are complementary — OPT-002 is the cheapest, most reversible option and the best-suited to genuine piloting given its weaker (GCC-inferred) evidence base, while OPT-003 is an already-evidenced, already-operating mechanism (ASM-016's inference problem does not apply to it) that DEC-009 and its own option record both recommend funding immediately, not gating behind a pilot. OPT-004 (advertising monetization) is the longer-horizon build. This plan's Three Horizons sequencing implements that resolved recommendation directly, not an assumed Egypt-first default or a re-derivation of the evidence tiering alone.

Three Horizons Roadmap

Horizon 1 (0-6 months) — Fund the near-term priority pair, continue Tier 1 at pace

- Continue at pace: OPT-001 (talabat mart densification) and OPT-005 (AI/personalisation)
- Pilot: OPT-002 (talabat pro acceleration in Egypt and Iraq) — DEC-009 names this the
- Fund now, within the existing bucket: OPT-003 (Food-leadership, UAE/Kuwait/Qatar) — per

Horizon 2 (6-18 months) — Evidence-gated decisions on the piloted/continuing items

Horizon 3 (18+ months) — Scale what H1/H2 evidenced, revisit the cross-market gap

Dependencies and Timelines

H2 decisions are explicitly gated on H1 evidence, not calendar time alone — per the KPI Tree's Governance family (KPI_Tree_v2.md), none of which currently has a baseline. The single largest dependency across all three horizons is the same one flagged in Section 2.2: no market-level allocation logic is disclosed anywhere in the corpus, so every horizon's country-specific decision (Egypt/Iraq for OPT-002, UAE/Kuwait/Qatar for OPT-003) is this plan's own proposed evidence-gate structure, not a validated talabat process.

Marketing Rollout Plan

Not primarily applicable to this plan's actual product (an internal capital-allocation decision-support system, per Section 5) — the "rollout" that matters is internal: to the extent H1's OPT-002 pilot involves customer-facing changes (accelerated pro marketing in Egypt/Iraq), that rollout should follow the same evidence-gated logic as the funding decision itself, not run ahead of it.


Team Expansion and Recruitment Goals

Not disclosed anywhere in the corpus for either the underlying AI capability or a capital-allocation function (Capital Allocation and Investment Governance.md's governance-mechanics gap extends here) — this plan does not invent a headcount figure. The genuine, evidenced need is process, not headcount: an owner for the KPI Tree's Governance family and the DEC-009/future-decision escalation path, a role this plan assigns functionally (Decision Steward) rather than as a proposed new hire.

Traceability summary (for Appendix/Section 14 use)

13. Monitoring and Evaluation

Capital release should follow leading proof signals before lagging financial outcomes



Source: KPI Tree v2; Section 13 | Classification: Selective KPI tree; full 50-KPI inventory retained in repository

Management implication: Do not scale options on spend alone; require adoption and operating proof before financial confirmation.

Exhibit 8. Capital release should follow leading proof signals before lagging financial outcomes.

Answer, stated first

This plan's monitoring framework does not invent new metrics to look complete — it reproduces KPI_Tree_v2.md's 50 KPIs across five families (Portfolio, Customer & growth, Financial, Operational, Governance) faithfully, states plainly which 30 already have a talabat-disclosed baseline and which 20 do not, and assigns no numeric target to any newly-instrumented KPI without a Decision Log entry to back it. The single most important finding this section carries forward: the entire Governance family (7 KPIs) has zero baseline, because talabat discloses zero internal capital-allocation governance mechanics anywhere in the 29-document corpus — any stage-gate or reallocation-threshold KPI this plan proposes is explicitly its own recommended mechanism.

KPIs Mapped to the Value Driver Tree

Full 50-KPI detail, leading/lagging tags, and geography tags: KPI_Tree_v2.md.

Leading vs. Lagging Split

Every KPI in KPI_Tree_v2.md carries this tag explicitly. The pattern that matters most for this plan's own decision logic: Portfolio and Operational KPIs are predominantly leading (spend pacing, dark-store density predict downstream GMV effect before it shows up in Financial KPIs); Financial KPIs are predominantly lagging (EBITDA margin confirms what already happened); Governance KPIs are a distinct third category — process-compliance indicators that are neither predictive of nor confirmatory about financial outcome, but are the precondition for the other two families being trustworthy inputs to a future allocation decision.

Churn — the Single Biggest Customer-Side Gap

No churn rate, definition, or cohort analysis exists anywhere in the corpus, for talabat overall (Group) or for Egypt specifically (Topics/Customer Churn.md) — the only country the corpus addresses this question for at all. This is carried forward from the pre-pivot problem's own finding, unchanged by the pivot — it becomes a metric this programme must start measuring from zero, not one it improves against a known baseline. Any churn target this plan or a future one proposes requires a Decision Log entry backing it; none exists yet.

Review Cadence

Root/Portfolio and Financial KPIs follow talabat's own quarterly reporting cadence (the corpus's actual disclosure rhythm — TLB-019, TLB-020, TLB-014), tracked against Section 9's base-case trajectory. The Governance family (7 KPIs, no baseline) and the newly-instrumented Customer & growth/Operational KPIs are reviewed at the same quarterly cadence but assessed trajectory-relative to the base case (sustained tracking toward or below it triggers escalation), not against an invented numeric threshold, per this section's own rule that no target may be set for a zero-baseline KPI without a Decision Log entry. This cadence is this OS's own proposed monitoring rhythm, anchored to talabat's disclosed reporting calendar — not a description of an internal talabat review process, none of which is disclosed (Capital Allocation and Investment Governance.md).

Governance KPIs — Explicitly the OS's Own Proposal, Not talabat's Process

Per Topics/Capital Allocation and Investment Governance.md's central finding, talabat discloses the programme's headline figures and one qualitative rationale, but no described approval committee, evaluation threshold, or stage-gate process. The seven Governance KPIs this plan proposes (e.g. "% of programme spend passing a documented stage-gate," "ROI/payback hurdle-rate compliance") are this OS's own recommended monitoring mechanism, stated as such — not a description or validation of talabat's actual internal process, and not to be cited elsewhere in this plan as if they were.

Tracking Tools and Feedback Loops

Not disclosed for talabat's own internal systems. This plan's own tracking mechanism is the vault-based architecture itself: Assumptions_Register.md (status transitions from Proposed to Approved), Investment_Options_Register.md and individual OPT- records (status transitions from candidate through approved/rejected/superseded), and the Decision Log's superseded-decision workflow (vault/Architecture/Decision_Management_Layer.md) as the feedback mechanism when new evidence changes a prior decision.

Kill/Pivot Criteria

Per Section 12's evidence-gated horizon structure: OPT-002 (the H1 pilot) is the clearest candidate for an explicit kill/pivot test — if Egypt/Iraq pro-adoption evidence at the H1→H2 gate falls materially short of the GCC+Jordan cohort's 20-28%/26-32% uplift range, the correct response is reassessment, not continued funding on the original ASM-016 inference. No numeric kill threshold is set here — per KPI_Tree_v2.md's own rule, that requires a dedicated Decision Log entry once H1 produces real Egypt/Iraq data, mirroring exactly how the pre-pivot plan's own DEC-007 handled the same class of problem (qualitative, trajectory-relative gates before real baseline data exists, not invented numeric thresholds).

Traceability summary (for Appendix/Section 14 use)

14. Appendices

Answer, stated first

This appendix is the plan's ghost deck: it does not introduce new arguments, it backs the ones the other thirteen sections already made. Three things live here and nowhere else in the plan: the full traceability note (§14.2), mapping

this plan's headline claims — the USD175mn/120mn/55mn envelope, the five ranked Investment Options, and each section's two-to-four most load-bearing numbers — to the vault note and underlying source document that supports each one; the DEC-009 allocation-range sensitivity exhibit (§14.3), the designated, and only, home for the full OPT-001–005 base/upside/downside cost table, shown here exactly once with its mandatory disclosure sentence and both of the Portfolio Register's own caveats, never as a headline financial commitment anywhere else in the plan; and an honest accounting of what this capstone project does and does not have — no team resumes exist (§14.4), no talabat-side pilot of the actual allocation recommendation has been run (§14.6), and no separate legal filings were produced (§14.7) — stated plainly rather than fabricated to make the appendix look more complete than the underlying project is.

14.1 Supporting Research and Data

The evidentiary infrastructure this plan draws on is not ad hoc — it is a maintained, auditable knowledge base with its own index and quality-control record:

- vault/MOC/Source Register.md — the document-level index for the 29-document primary corpus
- vault/MOC/Validation and Audit.md — every quality-control checkpoint run while building this

14.2 Traceability Note — Claim → Vault Note → Source

Per AI_Business_Plan_Template.md's explicit requirement and the drafting skill's own description of this section as "the closest to mechanically ready," the table below compiles this plan's actual argument — not every sentence in Sections 2–13, but every headline claim, the five Investment Options, the USD175mn/120mn/55mn figures, and each section's most load-bearing 2–4 numbers — into a single claim → vault note → source table. Geography tags are carried inline exactly as each source section states them, per Geographic_Evidence_Rules.md.

14.3 The DEC-009 Allocation-Range Sensitivity Exhibit (Designated Home)

Action title: talabat's disclosed 175mn/120mn/55mn envelope is fixed; this OS's own five-option cost ranges within it are an analytical recommendation, not a talabat allocation.

Option	Bucket	Base	Upside	Downside	Basis
OPT-001 — talabat mart dark-store densification	Everyday App	USD 45-55mn	USD 55-70mn	USD 30-40mn	ASM-029
OPT-002 — talabat pro acceleration (Egypt, Iraq)	Everyday App (opex)	USD 8-12mn	USD 12-20mn	USD 3-6mn	ASM-030
OPT-003 — Food-leadership CV P/partner-retention (UAE, Kuwait, Qatar)	Food-leadership	USD 25-32mn	USD 32-40mn	USD 15-22mn	ASM-031
OPT-004 — Advertising monetization gap closure	Enabling capability (cross-cutting, drawn from Everyday App opex)	USD 8-12mn	USD 12-18mn	USD 3-6mn	ASM-032
OPT-005 — AI/personalisation capability scaling	Enabling capability (cross-cutting, drawn from Everyday App opex)	USD 10-15mn	USD 15-22mn	USD 5-8mn	ASM-033

- Non-additive across the envelope. OPT-004 and OPT-005 are cross-cutting Enabling-capability options
- Partial coverage, not a full allocation. These five options do not claim to allocate 100% of either

14.4 Resumes of Key Team Members — Not Applicable

Stated honestly, not fabricated. Problem_Charter.md names six AASTMT MBA capstone team members (Group G02) with AASTMT student IDs, but every role assignment in that table is marked [assign] — no role has been finalized, and no biographical or professional-experience content for any team member exists anywhere in this project's corpus or vault. This plan is generated by the Agentic OS described throughout vault/Architecture/ (bp-orchestrator and seven specialist agent roles) rather than authored by named individual contributors with distinct functional biographies in the conventional business-plan sense. Inventing resumes to satisfy this template requirement would be exactly the fabrication CLAUDE.md's standing instructions and this plan's own evidence discipline prohibit — this section states the

gap instead. The team roster with its (unassigned) intended roles is reproduced from Problem_Charter.md for reference:

14.5 Technical Specifications of the AI Product

Full detail is Section 5's — this appendix gives the compressed specification, keeping the same explicit separation Section 5 established between the two distinct AI systems in scope:

- talabat's own embedded AI/personalisation capability (scaled, not built, by the 2026 investment
- This plan's own decision-support system (the Agentic OS) — a structured, multi-agent pipeline

14.6 Case Studies or Pilot Program Results

Two categories, kept explicitly separate — this section does not let either be mistaken for the other.

14.7 Legal and Compliance Documents

No separate legal or compliance filing was produced for, or by, this plan — this is a capstone Business Plan, not a regulatory submission, and no such document exists in this project's Outputs/ or elsewhere in the vault. What this section can honestly provide is a pointer to the disclosed compliance mechanics already compiled and cited in Sections 8 and 10, not a duplication of them:

- Inter-company/data-governance mechanism: the Delivery Hero Group Inter-Company Data Transfer
- A realized security/compliance incident: the December 2022 data breach (144,469 customers' personal

Traceability (for this section's own claims)

Document Control and Validation

14/14 sections independently verified; all five whole-plan gates PASS. Human approval remains mandatory before capital movement.