AI BUSINESS PLAN

talabat Group
AI-Enabled Investment Allocation · Version 1.1

How to allocate and govern the 2026 USD175 million programme across Everyday App and Food-leadership initiatives to maximize profitable growth and long-term platform economics.

Programme	AASTMT MBA · AI for Business Organizations
Team	Group G02
Scope	talabat Group · eight operating markets
Decision horizon	2026 investment programme
Document status	Version 1.1 enhancement · independently verified baseline · 25 July 2026

Recommendation in one line

Protect Food leadership in the evidenced GCC-3 markets, pilot talabat pro in Egypt then Iraq, continue proven Everyday App investments, and release funding through human-approved KPI gates.

Prepared by — Group G02

Name	Student Registration Number
Omar Hassan Ali Mohamadin Mohamed	24125108
Ahmed Fawzy Mohamed Rashed	24220998
Mohamed Magdy Abdullah Marawan	24224459
Sherif Samy Abdelhady Karoub	24225778
Ahmed Abd Elhakam Abd Elfattah Zaian	24224008
Abdallah Mahmoud Abdo Ahmed	24220430

Contents

1	Executive Summary
2	Business Description
3	Market Analysis
4	Value Proposition
5	AI Technology and Development
6	Business Model and Revenue Streams
7	Marketing and Sales Strategy
8	Operations Plan
9	Financial Plan
10	Risk Analysis
11	Corporate Social Responsibility & Responsible AI
12	Implementation Plan
13	Monitoring and Evaluation
14	Appendices

Evidence convention: disclosed facts are cited by TLB source/page or vault note; assumptions and forecasts retain ASM/node identifiers; Agentic OS allocations are explicitly labeled analytical.

1. Executive Summary

Approximately 69% of the disclosed 2026 programme is directed to Everyday App expansion

Source: TLB-020 pp.12,16; TLB-014 pp.16,19 | Classification: Disclosed talabat programme structure

Management implication: Govern release within the disclosed envelope; do not present the split as an Agentic OS allocation.

Exhibit 1. Approximately 69% of the disclosed 2026 programme is directed to Everyday App expansion.

The portfolio requires four actions—not one uniform funding decision

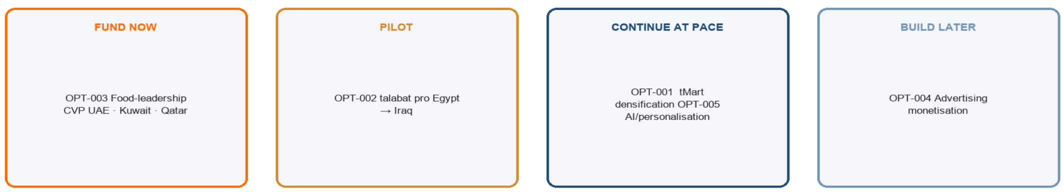

Source: Investment Options Register; DEC-009; OPT-001–005 | Classification: Agentic OS analytical recommendation; no dollar allocation shown

Management implication: Fund OPT-003 now, pilot OPT-002, continue OPT-001/005 at pace, and build OPT-004 later.

Exhibit 2. The portfolio requires four actions—not one uniform funding decision.

Situation

talabat Holding plc is a publicly listed, Group-wide on-demand delivery platform (Dubai Financial Market, IPO priced 29 November 2024) operating across eight markets — UAE, Kuwait, Qatar, Bahrain, Oman, Jordan, Iraq, and Egypt (Group; Problem_Charter.md).

Complication

talabat has not disclosed — and this project cannot observe — how that USD 175 million should be allocated within each bucket, across eight markets that differ sharply in maturity, profitability, and competitive intensity.

Question

How should talabat allocate its 2026 USD 175 million investment programme across Everyday App and Food-leadership initiatives — and within them, across markets and initiative categories — to maximise profitable growth, customer lifetime value, customer retention, and long-term platform economics, given the evidence…

Answer

This plan's answer is a framework and a committed sequence, not a single number. talabat does not need a new capital-allocation process invented from scratch — it needs the evidence-ranked, stage-gated discipline applied to the process it has already disclosed and already funded.

Business Name and Overview

talabat Holding plc — a mature, profitable, publicly listed Group operating a three-sided marketplace (customers, restaurant/retail Partners, riders) across eight countries, generating GMV of USD 9.5 billion in FY2025 across food delivery, grocery and quick-commerce (talabat mart), and a growing set of adjacent…

Mission and Vision

talabat's own stated trajectory — moving "from a multi-vertical food-delivery platform to the region's Everyday App," a framing management itself calls "a deliberate strategic choice, made from a position of strength" (Group; TLB-020, page 16) — is the mission and vision this plan's recommendation is built to serve, not a new one this plan proposes.

AI Product/Service, Target Market, and Value Proposition

The AI product is not customer-facing — it is an AI-enabled decision-support system that assembles talabat's own disclosed evidence, ranks candidate investment initiatives against explicit, documented criteria, builds range-bound (never single-point) financial scenarios, and monitors a KPI and stage-gate framework,…

Key Objectives, Growth Goals, and Financial/Funding Summary

Objective: raise the realized GMV, EBITDA, and customer-lifetime-value contribution of the already-approved USD 175 million programme by replacing undifferentiated spend with an evidence-ranked, stage-gated sequence, without requesting incremental capital beyond what talabat has already committed (Section_02_Business_Description.md §2.3; Section_09_Financial_Plan.md §9.6).

Traceability (full — this section synthesizes, rather than newly cites, Sections 2-14)

2. Business Description

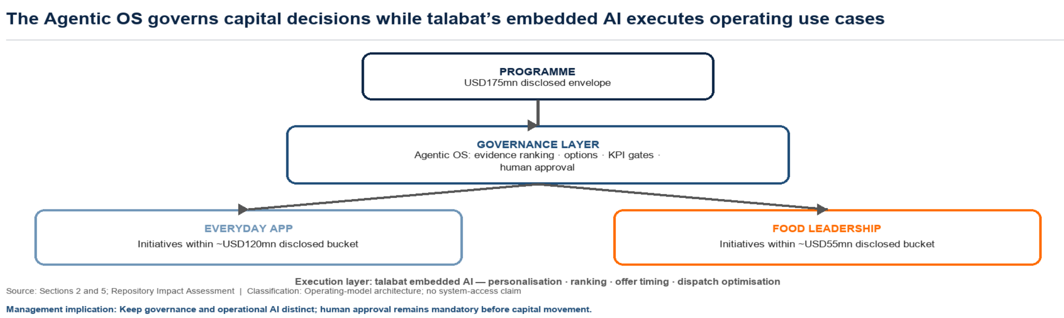

Exhibit 3. The Agentic OS governs capital decisions while talabat's embedded AI executes operating use cases.

Answer, stated first

talabat does not need to invent a new capital-allocation process — its Board has already approved and disclosed one: a USD175 million 2026 investment programme split into an Everyday App bucket (~USD120 million) and a Food-leadership bucket (~USD55 million), funded entirely from internal cash and framed explicitly as "investments rather than costs" (TLB-020, page 16).

2.1 Industry Overview and Opportunities for AI Adoption

talabat operates a three-sided marketplace (Strategic/How Talabat Creates Value.md; Facts/Marketplace_Facts.md) across eight countries (UAE, Kuwait, Qatar, Bahrain, Oman, Jordan, Iraq, Egypt — Entities/Countries.md), generating GMV of USD 7,428mn in FY2024 (+23% YoY) across its transacting verticals — food delivery,…

2.2 Problem Statement (MECE)

The business problem this plan addresses decomposes into four mutually exclusive, collectively exhaustive sub-problems, each corresponding to a genuine, separately-evidenced gap in what talabat discloses:

– Cross-bucket allocation (already resolved by talabat, not this plan's decision). The USD175mn

2.3 Governing Hypothesis

We believe that applying an AI-enabled, evidence-ranked, stage-gated allocation and monitoring framework to talabat's disclosed USD175 million 2026 investment programme — rather than leaving its within-bucket and cross-market allocation undifferentiated — will raise the programme's realized GMV, EBITDA, and…

2.4 Core Features and Benefits

The proposed AI-enabled decision-support system's core features, each already evidenced by a real artifact this analysis produced, not a hypothetical capability:

• Evidence assembly and geography discipline — every claim tagged Group/GCC/non-GCC/Egypt-standalone/

2.5 Business Model

The AI-enabled decision-support system is proposed as an internal capability, not a licensed or externally-sold product — consistent with talabat's own disclosed pattern of embedding AI/ML capability across functions rather than treating it as a discrete product line.

2.6 Current Business Stage

The underlying business (talabat) is public, mature, and profitable at Group level (IPO priced November 2024, TLB-025; FY2025 Adjusted EBITDA margin ~6.5% of GMV — Facts/Marketplace_Facts.md, TLB-002, pages 17 and 21) — not a startup.

Traceability (partial — full table due at Section 14)

3. Market Analysis

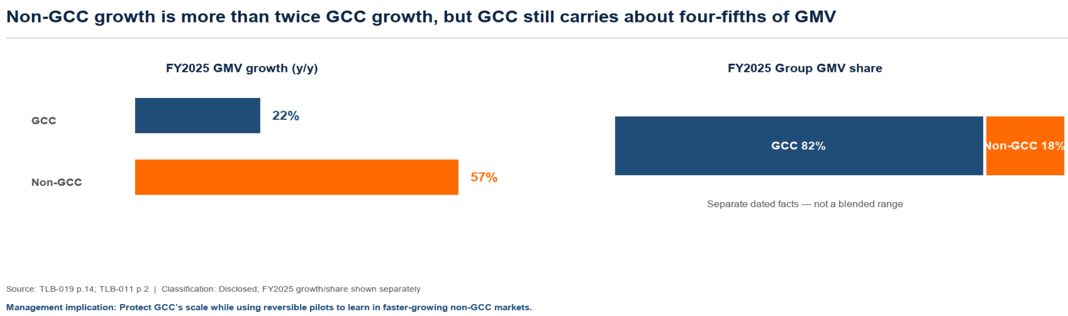

Exhibit 4. Non-GCC growth is more than twice GCC growth, but GCC still carries about four-fifths of GMV.

Answer, stated first

talabat's own disclosures show a company that already dominates the categories it operates in by category-share multiple, yet has penetrated only a fraction of the underlying addressable demand — and the 2026 investment programme's existing shape already reflects that gap more than it reflects the scale advantage.

3.1 Industry Trends — Market Sizing, Top-Down and Bottom-Up, Reconciled

Reframing note. The GSB template's Section 3 asks for "AI market size and growth projections" and "AI adoption trends... in the MENA region." Consistent with Section 2's framing (the AI-enabled capability this plan proposes is a capital-allocation decision-support layer, not a customer-facing product with its own…

3.2 Target Market — MECE Segmentation

Primary segmentation axis: talabat's own disclosed reportable-segment geography — GCC / non-GCC (Jordan + Iraq) / Egypt (standalone from FY2025). This is MECE on the IFRS 8 revenue-segment basis: every dollar of Group revenue and gross profit sits in exactly one of these three reportable segments, per talabat's own…

3.3 Competitive Analysis

SWOT — with a "so what" per quadrant

The four SWOT quadrants (internal/external × favourable/unfavourable) are MECE by construction; each is populated only with disclosed evidence, and each carries an explicit conclusion, not an orphan observation.

Strengths. Category leadership across all 8 markets (1x+–10x+ share vs. next-closest peer, TLB-001 p.5/TLB-002 p.5/TLB-014 p.4); an AI-personalised loyalty layer with a growing, management-estimated EBITDA contribution (>USD 14mn FY2024 → >USD 30mn FY2025 — worded by talabat as "estimated to generate/contribute," not a disclosed audited figure; Facts/AI_Facts.md; TLB-001, page 23; TLB-002, page 16); an AdTech monetization differentiator on the Partner side (an 8% CPG advertising investment ratio vs.

Weaknesses. An ~90–95% non-employed (third-party-logistics/freelance) rider workforce, which talabat itself flags as a labour-compliance and service-quality-control risk (TLB-001, pages 34, 47; TLB-026, page 138; provider-concentration detail at page 51); an AdTech monetization gap — Group-wide advertising revenue at 3.4–3.5% of GMV against management's own ~7%-of-GMV medium-term benchmark, a target reached only for talabat mart specifically in the UAE (Strategic/Competitive Weaknesses.md; TLB-001 p.21, TLB-002 pp.11/14/19, TLB-014 p.19); a…

Opportunities. Grocery/retail penetration of only ~1–3% against a USD 104–150bn+ addressable category (§3.1); the non-GCC geography's premium GMV growth rate (+57% y/y vs. GCC's +22%; GMV-by-geography basis, non-GCC incl. Egypt — TLB-019, p.14) while still representing only 18% (FY2025, TLB-019, p.14) to 21% (Q1 2026, TLB-011, p.2) of Group GMV — two dated data points, not a single blended figure; the disclosed AdTech monetization gap versus talabat's own ~7% benchmark; and the AI/personalisation capability's still-rising management-estimated EBITDA trajectory.

Threats. Competitive pressure has already measurably eroded non-high-value customer M1 retention by 4% y/y — but only in UAE, Kuwait, and Qatar specifically (TLB-019, page 9); FY2026 guidance names "a more competitive environment (new entrants and incumbents)" explicitly as a headwind (TLB-019, page 18); and Group Adjusted EBITDA margin compression is already realized, not merely guided — 6.0% (Q4 2025 baseline) stepping down to a 4.4–4.8% FY2026 guided range, with 4.8% already the actual Q1 2026 result (TLB-020, page 12 — the 6.0%→4.8% bridge — and page 14 — the 4.4–4.8% guidance; guidance origin TLB-019, pages 18–19) — the very margin step-down that funds this investment programme.

Explicit gap — Food-leadership's competitive-pressure evidence does not extend to Egypt or other non-GCC markets. The -4% non-high-value retention figure above is disclosed only for UAE, Kuwait, and Qatar (TLB-019, page 9; Topics/Food Leadership.md Open Questions). Applying this dynamic to Egypt or Jordan/Iraq would be an inferred-applicability extrapolation, not a directly evidenced parallel — this section does not make that extrapolation, and any future section that does must label it as such per Geographic_Evidence_Rules.md.

Competitor benchmarking

Talabat's own disclosures are notably thin on named rivals. Across the entire 29-document primary corpus, only one document — the International Offering Memorandum — names specific competitor brands: Deliveroo, Careem, noon, Jahez, and Snoonu (TLB-026, page 146; Entities/Competitors.md).

3.4 Opportunities

Untapped markets / underserved categories. Grocery/retail is the clearest untapped-opportunity category in the corpus: only ~1–3% penetrated of a USD 104–150bn+ addressable category (§3.1), and the vertical growing fastest by GMV (+47% y/y FY2025) despite carrying a structurally lower take rate/margin than Food — an…

Traceability (partial — full table due at Section 14)

4. Value Proposition

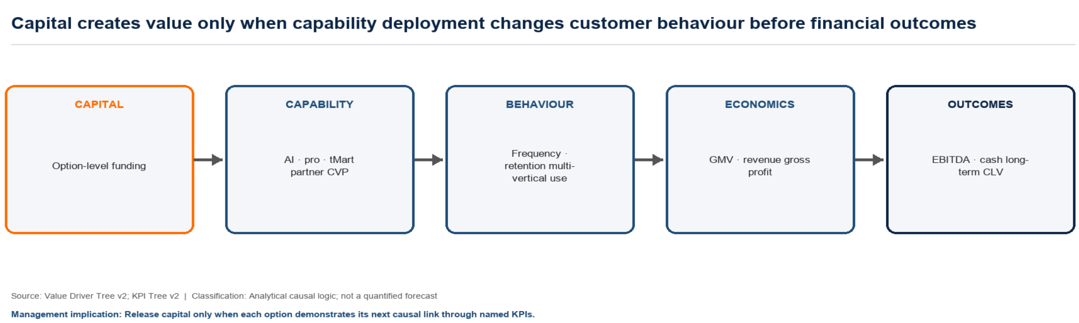

Exhibit 5. Capital creates value only when capability deployment changes customer behaviour before financial outcomes.

Answer, stated first

The value this plan creates does not come from a single new capability — it comes from applying a disciplined ranking-and-monitoring layer to five value mechanisms talabat has already proven — each at the geographic level it is actually disclosed — but has not yet evidenced as deliberately sequenced against its 2026 investment programme.

4.1 Unique Selling Proposition (USP)

The USP is not a new customer-facing product — it is evidence-ranked, stage-gated capital allocation applied to an already-committed USD175mn programme. talabat's own disclosed rationale for the programme is qualitative ("investments rather than costs...

4.2 Value Creation Mechanisms — Each Quantified

Why this breakdown is MECE: the seven mechanisms below are organized by the template's value categories — revenue lift (four distinct revenue mechanisms: multi-vertical engagement, subscription, targeted incentives, advertising monetization), cost savings/efficiency, risk reduction, and insight — mutually exclusive by value category and collectively exhaustive against the value drivers the corpus actually evidences.

4.3 Impact on Customer Outcomes and ROI — Calculation Logic Shown

Customer outcome, stated honestly: no absolute customer lifetime value figure is disclosed anywhere in the corpus (Topics/Customer Lifetime Value.md); the mechanisms above are evidenced as relative uplifts and GMV-share trends, not an absolute CLV number this section can cite.

Traceability (partial — full table due at Section 14)

5. AI Technology and Development

Answer, stated first

talabat does not need to build a new AI capability to execute this plan's core recommendation — the customer-facing AI capability (recommendation ranking, offer timing, dispatch optimisation) it will scale via the 2026 investment programme already exists as a hybrid: a logistics/dispatch stack licensed from parent…

5.1 Description of AI Technology

The two-capability split is MECE: talabat's embedded customer-facing AI and this plan's decision-support AI have distinct owners (talabat's product/tech organisation vs. this plan's governance process), distinct users (customers and riders vs.

5.2 Build / Buy / Partner Decision — Structured Options Analysis

talabat's embedded AI capability: a hybrid of partner and build — the logistics/dispatch layer (dispatch algorithms, real-time GPS tracking, rider staffing algorithm, picking tools) is licensed from Delivery Hero SE, the parent company (TLB-001, TLB-026, via Entities/Technology_Platforms.md), while the…

Criterion	Build (chosen — this OS)	Buy (third-party portfolio/allocation platform)	Partner (external analytics firm)
Cost	Low — assembled from existing agentic tooling over the project's own corpus; no licence or engagement fees	Higher — platform licence purchased for a scale and generality the five-option task does not require	Higher — recurring engagement fees for each allocation/monitoring cycle
Speed	Fast — operational within the capstone timeline; has already produced OPT-001–005, DEC-009, and the Forecast/KPI trees	Slower — procurement, onboarding, and corpus ingestion precede any first output	Slower — scoping and knowledge transfer precede any first output
Capability fit	High — purpose-built for evidence-tiered ranking with full citation traceability and geographic-evidence discipline, which is the task	Low-to-moderate — generic portfolio tooling; the citation-traceability and evidence-governance discipline would still have to be built around it	Moderate — strong generic analytics, but the evidence-governance conventions would live outside the process being governed
Data confidentiality	High — operates only on the project's corpus and internal decision records, inside the project boundary	Moderate — corpus, decision records, and working assumptions loaded into a third-party platform	Lower — decision rationale and working assumptions shared with an external firm

5.3 Proprietary Algorithms, Data Models

The corpus names one proprietary asset for talabat's embedded AI capability: a "Proprietary personalisation and ranking algorithm" — talabat's own ML-based recommendation system (TLB-001, page 23; Entities/Technology_Platforms.md).

5.4 Infrastructure and Tools

talabat's own infrastructure is not itemized in the corpus beyond the Delivery Hero licensing relationship. This plan's decision-support system runs on Claude Code / Cowork (documented in MEMORY.md's LLM decision log), with the vault-based knowledge/decision/forecast architecture described throughout vault/Architecture/.

5.5 Scalability and Adaptability

talabat pro — the subscription product whose offers the embedded AI capability times — is available in seven of the eight countries following the February 2025 Egypt launch, with Iraq the exception at that date (TLB-001, pages 11-12); at end-2024 the count was six of eight (TLB-015, page 77).

5.6 Research and Development Plans

For talabat's embedded AI: Strategic/AI Opportunities.md names talabat's own stated Group roadmap items (evidenced intent) — this plan does not add new customer-facing AI R&D beyond what talabat has already disclosed, since the programme's opex allocation is described as scaling, not inventing, capability.

Traceability (partial — full table due at Section 14)

6. Business Model and Revenue Streams

Answer, stated first

talabat's revenue is not four independent product lines but a single GMV-times-blended-take-rate system (~40–41% of GMV FY2024–FY2025; 38%→39% Q1 2025→Q1 2026) in which four disclosed fee-type components — commission fees, delivery & service fees, subscription fee & other income, and advertising & listing fees —…

6.1 Revenue Generation Methods — MECE Streams Traced to the Value Driver Tree

Why this is MECE. talabat's own "Management Revenue" reporting convention already decomposes total revenue into exactly four fee-type lines, net of one contra-revenue line — every disclosed revenue dollar belongs to exactly one of the five rows below, no overlap, no gap (Group; TLB-001, page 27; TLB-002, page 20):

Stream	FY2024	FY2025	Growth	Primary driver	Value Driver Tree node(s)
Commission fees	USD 1,062mn	USD 1,297mn	+25% → +22%	Food-vertical order volume (agent-model marketplace transaction — talabat earns commission without owning inventory); rate under some disclosed downward pressure	N-05, N-10 (Food-leadership capability), N-14/N-16 (order frequency), N-26 (Food GMV)
Delivery & Service fees	USD 696mn	USD 859mn	+29% → +24%	Order count and basket composition; structurally eroded per-order by talabat pro's free-delivery-above-threshold benefit even as pro raises order frequency	N-08 (pro adoption), N-14/N-16

Stream	FY2024	FY2025	Growth	Primary driver	Value Driver Tree node(s)
Subscription fee & Other Income (combined line — see decomposition below)	USD 952mn	USD 1,397mn	+44% → +47%	talabat pro subscriber-penetration (GMV share 32% (Q1 2025) → 49% (Q1 2026), Group; TLB-020 p.8; TLB-014 p.14) plus tMart's near-95%-take-rate owned-inventory "other income" component	N-07 (G&R product-mix shift), N-08, N-20/N-21/N-22 (retention/CLV uplift)
Advertising & listing fees	USD 246mn	USD 323mn	+27% → +32%	Partner demand for visibility — a lagging function of how engaged/frequently-ordering the customer base is, not an independent driver	N-09 (AI/personalisation), N-29, N-11/N-17 (multi-vertical usage), OPT-004
less: Vouchers and other discounts (contra-revenue)	—	USD (120)mn	—	Funded promotional/incentive spend	N-12 (targeted incentives)

6.2 Pricing Strategy

talabat's pricing is not one strategy but four distinct pricing mechanisms, one per revenue stream, each priced differently and each subject to a different constraint:

• Commission (agent-model, negotiated/rate-set) — talabat sets a commission rate on Food and Local

6.3 Recurring vs. One-Time Revenue Streams

Reframing note, consistent with Sections 2/3's own precedent. The template's language ("subscription fees, licensing, service contracts, custom AI solutions") describes a SaaS/professional-services revenue mix that does not map cleanly onto talabat's own disclosed marketplace model — talabat discloses no "one-time" revenue line in the traditional sense (no license sale, no setup fee, no custom-project contract).

• Contractually recurring: subscription fees only (talabat pro membership) — the smallest disclosed

6.4 Unit Economics — What One Customer and One Order Earn and Cost

Per-customer, per-month (the clearest disclosed unit-economics differentiator in the corpus). Food-only customers spent AED 194 (≈ USD 53)/month at 3.8 orders, while multi-vertical (Food + G&R) customers spent AED 814 (≈ USD 222)/month at 12.8 orders — a >4x monthly-spend gap tied to a ~3.4x order-frequency gap (Group, September 2024 basis; TLB-026, pages 91, 122).

6.5 The Proposed AI Decision-Support Layer's Own Revenue Treatment (cross-reference)

Consistent with Section 2.5's framing: the AI-enabled capital-allocation decision-support system this plan proposes is an internal capability, not a separately-sold product, and therefore does not add a new revenue line to the MECE breakdown above.

Traceability (partial — full table due at Section 14)

7. Marketing and Sales Strategy

Answer, stated first

talabat's own disclosed marketing posture is not price-led acquisition but customer-value-proposition (CVP) investment — management states explicitly that Food-leadership capital goes toward CVP "rather than matching competitor discounts and vouchers" (Group; TLB-020, page 16) — and this plan's marketing and sales…

7.1 Positioning — What talabat Is, and Deliberately Is Not, Competing On

The stated positioning. talabat frames itself as transitioning "from a multi-vertical food-delivery platform to the region's Everyday App" (Group; TLB-020, page 16, cited in Section_02_Business_Description.md §2.1), a narrative arc management itself traces across three successive brand statements — "Delivering food to…

7.2 Customer Acquisition — Channels, Funnel, and Journey Mapping

Disclosed channels (Fact, Group level). The corpus discloses acquisition and retention spend by cost category, not by discrete marketing channel (no disclosed breakdown of paid search, social, in-app referral, or offline spend exists anywhere in the corpus).

Funnel stage	Anchor to real, disclosed evidence	Illustrative conversion (stage-to-stage)
Awareness → App download / first visit	Not disclosed (no stated basis)	[SYNTHETIC — illustrative estimate, not derived from talabat's disclosed data]
Download → Trial (first order placed)	Not disclosed (no stated basis)	[SYNTHETIC, illustrative: ~25–35%]
Trial → Reaches ~6-order personalisation threshold	Real, disclosed mechanical threshold (Group; TLB-002, p.15) — but no disclosed % of trialists who reach it	[SYNTHETIC, illustrative: ~15–25% of first-time orderers]
6-order threshold → Multi-vertical adoption (Food + G&R)	Real, disclosed journey sequence (Group; TLB-013, p.6) — but no disclosed conversion %	[SYNTHETIC, illustrative: ~30–40%]
Multi-vertical → talabat pro sign-up	Real, disclosed journey terminus (Group; TLB-013, p.6) — but no disclosed conversion %	[SYNTHETIC, illustrative: ~10–20%]

7.3 Partnerships

Restaurant/retail Partners — the disclosed core of the partnership base. Food-leadership's stated partner-side investment is explicit: "on the partner side, we invest in retaining, winning back, and acquiring high-demand food partners...

7.4 Retention — Support, Community, and Product Enhancement

The ranked mechanism set, restated for the sales-and-marketing lens. Strategic/Customer Retention Drivers.md's own synthesis ranks five product-embedded retention levers by strength of corpus evidence: (1) talabat pro — a 20-28% order-frequency uplift and 26-32% retention uplift versus matched non-subscribers, and a…

7.5 Where This Positions the Plan's Own Ranked Options

Restated plainly: this section's positioning (CVP over discounting), acquisition approach (product- embedded journey stages plus a labeled synthetic funnel model, sequenced on the disclosed GCC-first-then- Egypt FinTech precedent), partnership base (restaurant/retail Partners, banking co-brand cards), and retention…

Traceability (partial — full table due at Section 14)

8. Operations Plan

Answer, stated first

talabat does not need to build new operational infrastructure to execute this plan's recommendation — its technology and delivery infrastructure is already licensed and operating at Group scale through a disclosed set of inter-company services agreements with parent Delivery Hero SE (the Global Licensing and Services Agreement and its category-specific variants), and its logistics/dispatch stack already runs across all eight markets.

8.1 Infrastructure Needs

The technology layer is inherited, not built. talabat's underlying infrastructure runs on capability licensed from Delivery Hero SE rather than owned or engineered standalone by talabat itself, delivered through a disclosed Global Licensing and Services Agreement (GLSA): members of the Group in Bahrain, Egypt, Iraq,…

8.2 Development and Maintenance Workflow

What is disclosed: a contractual services structure, not an internal software development lifecycle. No document in the corpus describes talabat's internal engineering process (sprint cadence, release process, environment/staging structure, or code-ownership model) — this is a genuine gap, stated openly rather than invented.

8.3 Key Team Roles and the Skills Gap vs. Today

Named roles the corpus discloses that this programme would plausibly draw on. talabat's executive roster includes several roles directly relevant to executing and governing this programme: Khaled Alfakesh, Chief Financial Officer since 2016, the most visible existing sponsor for a capital-allocation discipline…

8.4 7S Alignment Check

The corpus does not contain a pre-built 7S analysis for talabat — this is genuine synthesis, built here for the first time against the evidence assembled above and in the Topic Notes cited throughout this plan.

S	Rating	Evidence	Gap-closing action
Strategy	Ready	The Everyday App (~USD120mn) / Food-leadership (~USD55mn) split and the CVP-over-discounting positioning are clearly disclosed and consistently communicated (Section_02_Business_Description.md §2.1; Section_07_Marketing_and_Sales_Strategy.md §7.1; TLB-020, page 16)	None required — this is the plan's strongest evidenced "S"
Structure	Partially ready	Board dominated by Delivery Hero SE officers (Vandepitte/Chair also DH SE COO; Krause also DH SE General Counsel; Popp also DH SE interim CFO — Corporate Structure.md, TLB-026); Egypt operates through two dedicated legal entities (Delivery Hero Egypt SAE, Delivery Hero Dmart Egypt LLC) with no disclosed local board or decision rights (Corporate Structure.md Open Questions)	This plan's recommendation is presented for Group-level approval, consistent with the visible governance concentration, rather than assuming a country-level approval path exists (Decision-Making Process.md Strategic Implications)
Systems	Not ready — the largest gap	No internal capital-allocation committee, approval threshold, or stage-gate process for the USD175mn programme is disclosed anywhere in the corpus (Capital Allocation and Investment Governance.md)	This OS's own proposed mechanism, not a description of talabat's actual process: a lightweight stage-gate — each OPT- option reviewed against its own named stage gate (already specified per option, e.g. OPT-002's 2-quarter Egypt pilot checkpoint, OPT-001's margin-drag checkpoint) before scale-up funding, tracked against KPI_Tree_v2.md's Governance family (G1-G7, all seven newly-instrumented, none with a corpus baseline) — gated for human/team approval before any capital moves, per the Responsible-AI principle Section 11 develops further

S	Rating	Evidence	Gap-closing action
Shared Values	Partially evidenced	Management's own stated framing — "these are investments rather than costs" (TLB-020, page 16) and CVP-over-discounting (TLB-020, page 16) — is a real, quotable value statement; no equivalent explicit value statement about AI ethics or data governance is disclosed beyond the DTA's compliance mechanics (§8.5)	This plan's Section 11 (Responsible AI) is the designated place to state the value this OS itself operates on (human-approval-gated recommendations, never a false-precision automatic decision) explicitly, since talabat's own disclosure does not state an equivalent value for its internal allocation process
Skills	Not ready	No named data-science, capital-allocation, or portfolio-management function in the disclosed executive roster (§8.3); OPT-004/OPT-005 both name a specific, currently-unconfirmed capability gap (Partner-facing ad sales; per-market AI model tuning)	Locate the proposed Systems function (above) near the existing CFO/GRC structure rather than inventing a new seat; pilot OPT-005's per-market model-tuning gap in Egypt specifically, per OPT-005's own pilot recommendation
Style	Ready	A recognisable, repeated communication cadence — Capital Markets Day for medium-term targets, Annual Report for the year's strategic objective, quarterly results for dollar figures (Decision-Making Process.md "venue pattern") — and consistent language ("investments not costs," CVP-over-discounting) across TLB-014/019/020	None required
Staff	Partially ready	Governance-adjacent named roles exist (GRC, Internal Audit, Legal/Board Secretary — §8.3) — but two leadership transitions are recorded in the corpus without any narrated explanation: CEO Tomaso Rodriguez → Toon Gyssels (between TLB-018, Aug 2025, and TLB-009, Feb 2026) and board seat Muhammad Hussain Ghati Al Jbori → Abdul Wahab Al-Halabi (between TLB-004/TLB-026 and TLB-008) (Corporate Structure.md, Entities/Executives.md)	Named explicitly here as a continuity caveat and carried forward into Section 10 (Risk Analysis) as an organizational risk, not resolved or explained by this plan, since no source in the corpus does so either

8.5 Security and Compliance Measures

A disclosed security incident, not a hypothetical risk. In December 2022, talabat was hacked by an external attacker based in Norway, who gained access to the personal data of 144,469 customers in one of talabat's markets (market unnamed in the disclosure); talabat informed the competent data-protection regulator, which opened an investigation, and talabat paid a USD 150,000 penalty (Group, specific market not disclosed; TLB-026, page 49).

Traceability (partial — full table due at Section 14)

9. Financial Plan

Answer, stated first

talabat's own FY2026 guidance already prices in the 2026 investment programme's cost: GMV growth of 11-14% cFX, Revenue growth 14-17% cFX, Adjusted EBITDA of 4.4-4.8% of GMV (down from 6.0% in Q4 2025), and Free Cash Flow of 3.2-3.6% of GMV.

9.1 Value Driver Tree — Linking the Investment Programme to Financial Outcomes

The governing chain (vault/Knowledge/Investment_Relationship_Map.md; formalized as Value_Driver_Tree_v2.md's 45 nodes): Investment → capability deployment → adoption/operational change → customer/partner behaviour → order frequency → multi-vertical usage → basket/AOV → retention/CLV → GMV → revenue → gross profit → EBITDA → cash flow.

9.2 Initial Investment and Operational Costs

Disclosed, not decided by this plan: USD175mn total, Board-approved February 2026, fully funded by internal cash — ~USD120mn Everyday App (~USD75mn opex + ~USD45mn capex) + ~USD55mn Food-leadership (vault/Decisions/Investment_Portfolio_Register.md).

9.3 Revenue Forecasts — Base, Upside, Downside (FY2026)

Scenario	GMV growth (cFX)	Revenue growth (cFX)	Adj. EBITDA margin	FCF margin	Basis
Base	11-14%	14-17%	4.4-4.8% of GMV	3.2-3.6% of GMV	talabat's own disclosed FY2026 guidance (TLB-020, TLB-014)
Upside	Toward/above 14%	Toward/above 17%	Toward 4.8%+	Toward/above 3.6%	Everyday App/Food -leadership mechanisms (multi-vertical, pro, advertising) compound faster than guidance embeds — Scenarios_v2.md upside case
Downside	Toward/below 11%	Toward/below 14%	Toward/below 4.4% (Q1 2026 actual already 4.8%, -9% y/y in absolute EBITDA dollars)	Toward/below 3.2%	Margin compression deepens, behaviour-change trends moderate faster than guidance assumes — Scenarios_v2.md downside case

All three rows are Group-level, per DEC-008. No bucket-level, market-level, or initiative-level breakout of these figures is shown here — none is disclosed, and constructing one would be exactly the false-precision error DEC-008 was escalated to prevent.

9.4 Break-Even Analysis and P&L Snapshot — What Can and Cannot Be Shown

A conventional break-even/payback analysis requires a return figure to compare against the USD175mn outlay. No such figure exists in the corpus for either bucket or any candidate initiative — confirmed independently by Investment_Relationship_Map.md, EBITDA.md, and Capital Allocation and Investment Governance.md.

9.5 Key Assumptions Register (the numbers this case depends on)

All fourteen listed are status: Approved in vault/Decisions/Assumptions_Register.md (resolved 2026-07-23, alongside DEC-009) — citable as settled inputs per the register's own rule. Listed here for transparency, exactly as the pipeline requires:

ID	Statement (compressed)	Confidence
ASM-015	Basket/AOV moves directionally with multi-vertical GMV share (no disclosed AOV figure exists)	Low

ID	Statement (compressed)	Confidence
ASM-016	GCC+Jordan pro/multi-vertical uplift statistics applied to Egypt/non-GCC as inference, not fact	Low
ASM-017	FY2026 GMV guidance illustratively split into baseline vs. programme-attributable components	Low
ASM-018	~2.2:1 Everyday App:Food-leadership GMV-effect split — restricted per DEC-008, narrative only	Low
ASM-019	FY2026 margin step-down sees directional partial recovery in FY2027+, timing/magnitude unstated	Low
ASM-020	Egypt's ≈USD19.3m pro-rata programme exposure — illustrative worked-example only, not a budget line	Low
ASM-021–027	Module-level continuation-trend projections (pro, multi-vertical, AI, incentives, advertising, G&R, Food)	Medium
ASM-028	Logistics improvement continues, directional-only, no numeric FY2026 target	Low

9.6 Funding Requirements, Allocation, and Monetization Strategy

Funding requirement: none beyond what talabat has already committed — the USD175mn programme is fully funded from internal cash (TLB-020, page 16); this plan does not request incremental capital beyond the disclosed envelope.

Traceability (partial — full table due at Section 14)

10. Risk Analysis

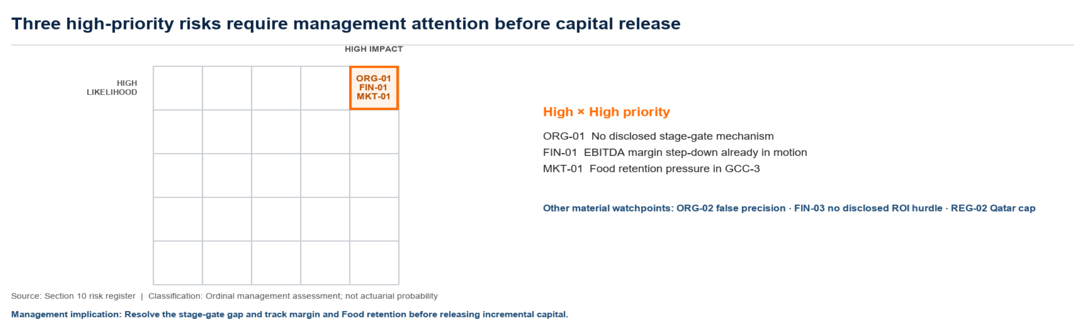

Exhibit 7. Three high-priority risks require management attention before capital release.

Answer, stated first

This plan's largest risk is not algorithmic — it is that talabat is funding a margin-dilutive, evidence-thin capital-allocation programme with almost no disclosed internal mechanism to correct course, layered on top of already-realized competitive, regulatory, and leadership volatility.

10.1 Technical Risks

TECH-01 — AI/personalisation effectiveness is unconfirmed outside the GCC+Jordan-measured cohort. talabat's AI/ML personalisation stack drives a disclosed, multi-year EBITDA contribution (USD14mn+ FY2024 → USD30mn+ FY2025, Group; TLB-002, p.15, per Topics/AI.md), but the corpus does not disclose whether models are tuned separately for non-GCC markets or run on the same Delivery Hero-inherited regional models (AI.md Open Questions).

10.2 Market Risks

MKT-01 — Competitive pressure is already eroding Food's non-high-value customer retention faster than its capital allocation defends against it. The disclosed non-high-value M1 retention decline is -4% y/y, scoped specifically to UAE, Kuwait, and Qatar (market-comparison: UAE, Kuwait, Qatar; TLB-019, p.9) —…

10.3 Financial Risks

FIN-01 — The disclosed FY2026 Adjusted EBITDA margin step-down is real, evidenced, and already in motion, with a documented trail — not a hypothetical risk. Group Adjusted EBITDA margin fell from 6.7% (FY2024) to 6.5% (FY2025); FY2026 guidance steps this down further to 4.4-4.8% of GMV, a deliberate choice funding the USD175mn programme, not a surprise (Group; TLB-020; Strategic/Strategic Risks.md).

10.4 Organizational & Governance Risks

ORG-01 — No internal capital-allocation committee, approval threshold, or stage-gate process is disclosed anywhere in the corpus. This is the single largest, load-bearing governance gap this plan works around: Topics/Capital Allocation and Investment Governance.md states plainly that no document describes how a…

10.5 Regulatory Risks

REG-01 — Antitrust/competition-law exposure tied directly to talabat's market dominance. Ongoing Kuwait Competition Protection Authority investigations/litigation over Partner dealings, UAE Competition and Consumer Protection Department notices over Partner-subscription contract clauses, and an Iraq…

10.6 Probability-Impact Matrix

New synthesis, built fresh against the evidence assembled above — not a pre-existing corpus artifact. Probability and Impact are each rated High/Medium/Low based on: Probability — whether the risk is already observed/realized in the corpus (High), a plausible extension of an observed trend or named but unresolved…

Probability ↓ / Impact →	Low	Medium	High
High	—	FIN-03, REG-02	FIN-01, MKT-01, ORG-01
Medium	—	TECH-01, MKT-02, MKT-03, FIN-02, FIN-04, ORG-03, REG-01, REG-04	TECH-02, ORG-02
Low	REG-03	MKT-04	—

10.7 Pre-Mortem: "It is FY2028, and the 2026 investment programme failed to deliver a more loyal

– **The governance gap never closed, and the plan's own numbers were mistaken for commitments

– **Margin compression deepened past the point of return, and no one could tell which line was responsible

10.8 Mitigation Strategy and Owners (per material risk)

Owners are drawn from talabat's disclosed executive roster (Entities/Executives.md) where a function plausibly exists; where no such function is named — a gap this plan states explicitly rather than inventing a role — the nearest disclosed function is named as interim owner, consistent with Section_08_Operations_…

ID	Risk	Prob.	Impact	Mitigation	Owner
TECH-01	AI/personalisation effectiveness unconfirmed outside GCC+Jordan	Med	Med	Pilot OPT-005's model-tuning specifically in Egypt before crediting Group-level EBITDA trends to non-GCC markets; track via Operational/Governance KPIs, not assumed parity with GCC	Pedram Assadi, COO (no CDO/Head of Data Science is named in the corpus — a disclosed skills gap, §8.3)

ID	Risk	Prob.	Impact	Mitigation	Owner
TECH-02	Parent-infrastructure security dependency (Dec 2022 breach)	Med	High	Continue relying on the disclosed DTA/Binding Corporate Rules/GDPR-SCC framework as the only disclosed cross-entity control; monitor incident recurrence and regulator engagement explicitly	Mohamd Abu Amara (Head of GRC); Abdullah AlGhrawi (VP Legal, GRC & Board Secretary)
MKT-01	Competitive pressure vs. Food-leadership's ~USD55mn defense budget	High	High	Hold CVP investment to its stated CVP-over-discounting logic (TLB-020, p.16); monitor the -4% y/y M1 retention metric quarterly; define and deploy the "dry powder" reserve's trigger condition explicitly	Pedram Assadi, COO (no single named Food-vertical commercial owner is disclosed)
MKT-02	Rider/delivery-workforce dependency (~90-95% 3PL/freelance)	Med	Med	Treat fulfilment capacity as a monitored dependency, not an assumed input, for any option (OPT-001, OPT-002) that increases order volume; add a fulfilment-capacity check to the ORG-01 stage gate	Pedram Assadi, COO
MKT-03	AdTech monetization gap not closing at Group scale	Med	Med	Fund OPT-004 as a longer-horizon build per DEC-009's tiering, not an assumed automatic transfer of the UAE benchmark; track AdTech revenue as % of GMV by market	Wassim Makarem, SVP Grocery & Retail (no dedicated AdTech/commercial-advertising executive is named)
MKT-04	Recurrence of geopolitical/operational disruption	Low	Med	Monitor order-volume recovery patterns by market as a leading indicator rather than assuming management's "transitory" framing holds without independent evidence	Pedram Assadi, COO; Board (Pieter-Jan Vandepitte, Chair) for Group-wide escalation

ID	Risk	Prob.	Impact	Mitigation	Owner
FIN-01	FY2026 Adjusted EBITDA margin step-down deepening beyond guidance	High	High	Track quarterly margin against the FY2026 4.4-4.8% guided range via KPI G6; treat two consecutive quarters below 4.4% as the ORG-01 stage-gate review trigger	Khaled Alfakesh, CFO
FIN-02	Egypt FX/currency exposure, unquantified magnitude	Med	Med	Model any Egypt-specific illustration (e.g. ASM-020's ≈USD19.3mn pro-rata proxy) as FX-sensitive, consistent with management's "GMV inflation-hedge" framing; do not attempt to quantify the depreciation itself	Khaled Alfakesh, CFO
FIN-03	No disclosed ROI/payback/hurdle-rate criterion for either bucket	High	Med	Adopt Governance KPI G5 (ROI/payback-period hurdle-rate compliance rate) as the recommended tracking mechanism; until adopted, present every allocation recommendation as a range or staged proposal, never an implied-return claim	Khaled Alfakesh, CFO, jointly with the proposed Systems function (ORG-01)
FIN-04	G&R margin drag widening in lower-maturity non-GCC/Egypt markets	Med	Med	Track G&R EBITDA margin drag specifically as store density increases in non-GCC/Egypt; treat OPT-001's own margin-drag checkpoint as the relevant stage gate	Wassim Makarem, SVP Grocery & Retail
ORG-01	No disclosed internal capital-allocation stage-gate mechanism	High	High	Adopt the stage-gate mechanism proposed in Section _08_Operations_Plan.md §8.4 — each OPT- option reviewed against its own named stage gate before scale-up funding, tracked against KPI family G1-G7 — explicitly labeled as this OS's own proposal pending human/team adoption	Proposed Systems function, located near CFO/GRC (Alfakesh; Abu Amara; AlGhrawi) — no equivalent function is currently named

ID	Risk	Prob.	Impact	Mitigation	Owner
ORG-02	False-precision risk on this OS's own constructed allocation ranges	Med	High	Every constructed range (ASM-029–033) carries the mandatory DEC-009 disclosure sentence wherever shown; appears only as a labeled sensitivity/appendix exhibit, never a headline exhibit (DEC-008/DEC-009); requires explicit human/team sign-off before any capital reallocation executes	decision-steward (this OS), with mandatory human/team (Business Plan Lead / Board-level) sign-off
ORG-03	Leadership continuity — two unexplained recent transitions	Med	Med	Name this risk explicitly to the Board/sponsor at proposal time; anchor accountability for this plan's recommendations to named functions (CFO, GRC cluster), not to the two most recently vacated seats specifically	Board (Pieter-Jan Vandepitte, Chair)
REG-01	Antitrust/competition-law exposure (Kuwait, UAE, Iraq)	Med	Med	Factor active antitrust exposure into commission-economics assumptions for any option touching these markets (OPT-003's UAE component); monitor, since no resolution timeline is disclosed for any of the three inquiries	Abdullah AlGhrawi, VP Legal, GRC & Board Secretary
REG-02	Realized commission-rate cap in Qatar	High	Med	Model OPT-003's Qatar-specific commission economics against the already-realized MOCI cap, not the uncapped GCC average	Abdullah AlGhrawi, VP Legal, GRC & Board Secretary; Wassim Makarem for commercial-model impact
REG-03	Oman licensing dispute	Low	Low	Monitor for resolution/escalation; no OPT- option or dollar figure currently ties to Oman by name	Abdullah AlGhrawi, VP Legal, GRC & Board Secretary

ID	Risk	Prob.	Impact	Mitigation	Owner
REG-04	Data-protection regulatory exposure; no jurisdiction-specific statute disclosed for any of 8 markets	Med	Med	Continue operating this plan's decision-support layer on Group/GCC/segment-level aggregates only (no individual customer personal data processed, §8.5); state the jurisdiction-specific-law gap explicitly rather than inventing citations	Mohamd Abu Amara, Head of Governance, Risk & Compliance

Traceability (partial — full table due at Section 14)

11. Corporate Social Responsibility & Responsible AI

Answer, stated first

This plan treats Responsible AI as risk management, not decoration, because it has to: this OS's own recommendation process creates a real risk (ORG-02, the false-precision risk DEC-008 and DEC-009 were each escalated to guard against), sits inside an organization with no disclosed correction mechanism for its own…

11.1 Ethical AI Commitments Tied to Concrete Controls

This breakdown is MECE by control type, not by ethical principle in the abstract: each commitment names exactly one operating control, the Section 10 risk (or, for Control 4, the stakeholder expectation) it answers, and the evidence that the control already works in practice, not just in policy language.

Commitment	Concrete control	Section 10 risk / stakeholder expectation	Evidence the control is real, not aspirational
No false-precision capital numbers	Headline-exhibit restriction + mandatory disclosure sentence on every constructed range	ORG-02	DEC-008 (Option 1, approved); DEC-009 (middle path, approved)
Geography-crossing inferences always labeled	Nine-tag geography system; checked at citation audit (13.9) and Geographic Evidence Gate (Stage 16)	Fairness/evidentiary-integrity risk this pivot's own root-cause analysis identified	Geographic_Evidence_Rules.md; the superseded Value_Driver_Tree.md case study
No customer-level data claim	Aggregate-only inputs (Group/GCC/segment)	Stakeholder expectation (Board/investor trust, post-TECH-02/REG-04)	Problem_Charter.md; Section_08_Operations_Plan.md §8.5
No assumed-uniform AI effectiveness across markets	OPT-005 per-market model-tuning pilot, Egypt first	TECH-01	Section_08_Operations_Plan.md §8.4; Section_10_Risk_Analysis.md §10.8

11.2 Sustainability Initiatives

talabat discloses real, if partial, sustainability activity — this plan reports it honestly, including what it does not know. Four categories of disclosed activity exist in the primary corpus:

• Electric-vehicle delivery fleet. talabat operates 240 electric vehicles in the UAE (as of January

11.3 Commitment to Responsible AI Practices and Governance Structures

The governing principle, stated at charter level and never diluted downstream. Problem_Charter.md states: "The AI does not claim to possess talabat's internal, customer-level data, and its output is not a substitute for management decision-making...

11.4 What this section does not claim

Consistent with the discipline the rest of this plan applies, this section states three things it explicitly does not do: (1) it does not claim talabat has adopted, or has committed to adopt, any of the governance structures proposed above — they are this OS's own recommendation, pending human/team decision, exactly…

Traceability (partial — full table due at Section 14)

12. Implementation Plan

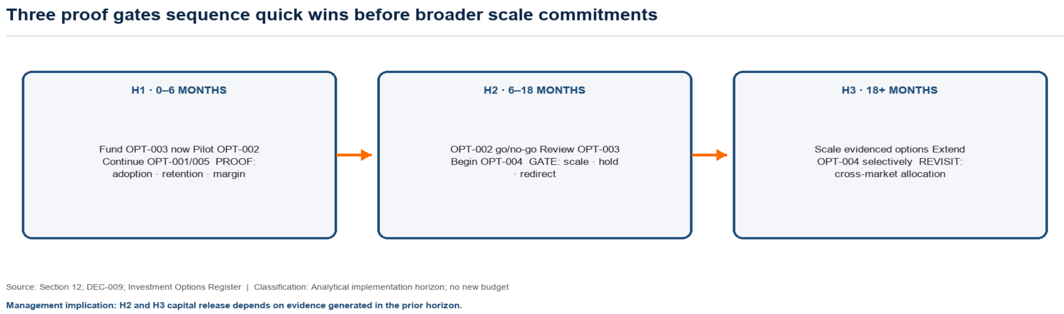

Exhibit 6. Three proof gates sequence quick wins before broader scale commitments.

Answer, stated first

The five candidate Investment Options this plan has identified do not all belong at the same starting point. DEC-009's underlying 11-criterion comparison groups OPT-001 (talabat mart densification) and OPT-005 (AI/personalisation scaling) as the strongest-evidence, broadest-applicability pair — continue-at-pace items requiring no new pilot.

Three Horizons Roadmap

Horizon 1 (0-6 months) — Fund the near-term priority pair, continue Tier 1 at pace

• Continue at pace: OPT-001 (talabat mart densification) and OPT-005 (AI/personalisation

• Pilot: OPT-002 (talabat pro acceleration in Egypt and Iraq) — DEC-009 names this the

Horizon 2 (6-18 months) — Evidence-gated decisions on the piloted/continuing items

• OPT-002 go/no-go: contingent on H1's Egypt/Iraq pro-adoption evidence. If frequency/retention

• OPT-003 review: using H1's budget-weighting and dry-powder-trigger baseline, assess whether the

Horizon 3 (18+ months) — Scale what H1/H2 evidenced, revisit the cross-market gap

• OPT-002: scale if H2's go/no-go supported it; hold or reassess if it did not clear its evidence

• Revisit the largest remaining open question from Section 2.2 (problem 3: cross-market allocation) —

Dependencies and Timelines

H2 decisions are explicitly gated on H1 evidence, not calendar time alone — per the KPI Tree's Governance family (KPI_Tree_v2.md), none of which currently has a baseline. The single largest dependency across all three horizons is the same one flagged in Section 2.2: no market-level allocation logic is disclosed…

Marketing Rollout Plan

Not primarily applicable to this plan's actual product (an internal capital-allocation decision-support system, per Section 5) — the "rollout" that matters is internal: to the extent H1's OPT-002 pilot involves customer-facing changes (accelerated pro marketing in Egypt/Iraq), that rollout should follow the same…

Team Expansion and Recruitment Goals

Not disclosed anywhere in the corpus for either the underlying AI capability or a capital-allocation function (Capital Allocation and Investment Governance.md's governance-mechanics gap extends here) — this plan does not invent a headcount figure.

Traceability summary (for Appendix/Section 14 use)

13. Monitoring and Evaluation

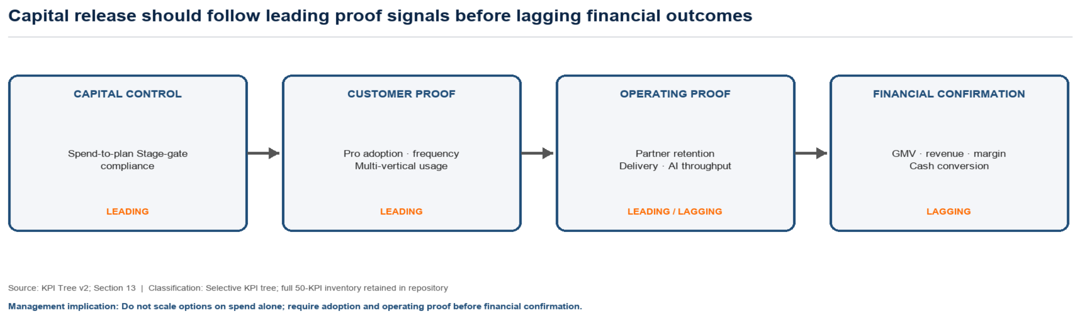

Exhibit 8. Capital release should follow leading proof signals before lagging financial outcomes.

Answer, stated first

This plan's monitoring framework does not invent new metrics to look complete — it reproduces KPI_Tree_v2.md's 50 KPIs across five families (Portfolio, Customer & growth, Financial, Operational, Governance) faithfully, states plainly which 30 already have a talabat-disclosed baseline and which 20 do not, and assigns no numeric target to any newly-instrumented KPI without a Decision Log entry to back it.

KPIs Mapped to the Value Driver Tree

Family	Count	Baseline-exists	Newly-instrumented	Representative KPIs
Portfolio	8	4	4	Spend-to-plan by bucket (baseline: disclosed EBITDA bridge); % of programme spend passing a documented stage-gate (new)
Customer & growth	12	9	3	talabat pro subscriber GMV share (baseline: 32%→49%, Q1 2025→Q1 2026); multi-vertical GMV share (baseline: 68%→73%→76%, Dec'24→Dec'25→Mar'26); churn rate (new — see below)
Financial	14	11	3	Adj. EBITDA margin (baseline: guided range); programme-attributable incremental GMV (new, ASM-017-derived, illustrative only)
Operational	9	6	3	Dark-store count/density (baseline: 130→160 stores); delivery time/on-time rate (baseline: disclosed trend)
Governance	7	0	7	% of spend passing a documented stage-gate; ROI/payback hurdle-rate compliance — entire family newly-instrumented

Full 50-KPI detail, leading/lagging tags, and geography tags: KPI_Tree_v2.md.

Leading vs. Lagging Split

Every KPI in KPI_Tree_v2.md carries this tag explicitly. The pattern that matters most for this plan's own decision logic: Portfolio and Operational KPIs are predominantly leading (spend pacing, dark-store density predict downstream GMV effect before it shows up in Financial KPIs); Financial KPIs are predominantly…

Churn — the Single Biggest Customer-Side Gap

No churn rate, definition, or cohort analysis exists anywhere in the corpus, for talabat overall (Group) or for Egypt specifically (Topics/Customer Churn.md) — the only country the corpus addresses this question for at all.

Review Cadence

Root/Portfolio and Financial KPIs follow talabat's own quarterly reporting cadence (the corpus's actual disclosure rhythm — TLB-019, TLB-020, TLB-014), tracked against Section 9's base-case trajectory.

Governance KPIs — Explicitly the OS's Own Proposal, Not talabat's Process

Per Topics/Capital Allocation and Investment Governance.md's central finding, talabat discloses the programme's headline figures and one qualitative rationale, but no described approval committee, evaluation threshold, or stage-gate process.

Tracking Tools and Feedback Loops

Not disclosed for talabat's own internal systems. This plan's own tracking mechanism is the vault-based architecture itself: Assumptions_Register.md (status transitions from Proposed to Approved), Investment_Options_Register.md and individual OPT- records (status transitions from candidate through…

Kill/Pivot Criteria

Per Section 12's evidence-gated horizon structure: OPT-002 (the H1 pilot) is the clearest candidate for an explicit kill/pivot test — if Egypt/Iraq pro-adoption evidence at the H1→H2 gate falls materially short of the GCC+Jordan cohort's 20-28%/26-32% uplift range, the correct response is reassessment, not continued funding on the original ASM-016 inference.

Traceability summary (for Appendix/Section 14 use)

14. Appendices

Answer, stated first

This appendix is the plan's ghost deck: it does not introduce new arguments, it backs the ones the other thirteen sections already made. Three things live here and nowhere else in the plan: the full traceability note (§14.2), mapping this plan's headline claims — the USD175mn/120mn/55mn envelope, the five ranked…

14.1 Supporting Research and Data

The evidentiary infrastructure this plan draws on is not ad hoc — it is a maintained, auditable knowledge base with its own index and quality-control record:

- vault/MOC/Source Register.md — the document-level index for the 29-document primary corpus

14.2 Traceability Note — Claim → Vault Note → Source

Per AI_Business_Plan_Template.md's explicit requirement and the drafting skill's own description of this section as "the closest to mechanically ready," the table below compiles this plan's actual argument — not every sentence in Sections 2–13, but every headline claim, the five Investment Options, the USD175mn/120mn/55mn figures, and each section's most load-bearing 2–4 numbers — into a single claim → vault note → source table.

Sec.	Claim	Vault note	Source
2	USD175mn = ~120mn Everyday App + ~55mn Food-leadership (Group, Board-approved Feb 2026, internally funded)	Investment_Portfolio_Register.md	TLB-020, pp.12,16; TLB-014, pp.6,19
2	"Investments rather than costs... expected to offset the margin impact over time"	—	TLB-020, page 16
2	No disclosed internal capital-allocation governance mechanics (committee, approval threshold, stage-gate)	Topics/Capital Allocation and Investment Governance.md	corpus-wide absence finding
2	Five candidate Investment Options identified and ranked	Investment_Options_Register.md; OPT-001–005	DEC-009
2	FY2024 GMV USD7,428mn (+23% YoY), Group	Topics/GMV.md	TLB-001, pages 6, 26
2	Eight operating markets (UAE, Kuwait, Qatar, Bahrain, Oman, Jordan, Iraq, Egypt)	Entities/Countries.md	Problem_Charter.md
2	Dubai Financial Market listing, November 2024	—	TLB-025
2	FY2025 Adjusted EBITDA margin ~6.5% of GMV, Group	Facts/Marketplace_Facts.md	corpus-wide
3	FY2025 GMV USD9.5bn (+28% y/y cFX), Group; ~7.7mn customers, ~585mn orders, ~84k Partners	—	TLB-019, pp.5,13; TLB-002, page 12
3	Category share 1x+–10x+ across 8 markets vs. next-closest peer	—	TLB-001 p.5; TLB-002 p.5; TLB-014 p.4
3	Grocery/retail penetration ~1–3% of a USD104–150bn+ addressable category (two unreconciled TAM figures, both stated)	—	TLB-014, pages 5, 19; TLB-020, page 16
3	GCC/non-GCC/Egypt FY2025 segment revenue, gross profit, net profit (Egypt-standalone; GCC/non-GCC country-comparison)	Topics/GCC vs non-GCC.md	TLB-002, page 111
3	Non-GCC GMV +57% y/y vs. GCC +22% y/y, FY2025 (market-comparison)	—	TLB-019, page 14
3	talabat pro subscriber GMV share 32%→49% (Q1 2025→Q1 2026), Group	Topics/Talabat Pro.md	corpus-wide
3	Multi-vertical GMV share 68%→73%→76% (Dec'24→Dec'25→Mar'26), Group	Topics/Multi-Verticality.md	TLB-019 p.7; TLB-020 p.8
3	Egypt category-share figure inconsistent across the corpus (10x+ vs. 1x+→4x+), both stated, neither resolved	Topics/Egypt.md; Topics/Competition.md	TLB-001 p.5; TLB-002 p.5; TLB-014 p.4; TLB-026 p.114
3	Non-high-value M1 retention -4% y/y (market-comparison: UAE, Kuwait, Qatar only)	Topics/Food Leadership.md	TLB-019, page 9
3	Only Deliveroo/Careem/noon/Jahez/Snoonu named anywhere in the primary corpus, all from one document	Entities/Competitors.md	TLB-026, page 146

Sec.	Claim	Vault note	Source
4	Multi-vertical 13.0x vs. 3.8x order frequency (Group)	Topics/Multi-Verticality.md; Value_Driver_Tree_v2.md N-14 (order frequency), N-11 (GMV-share trend)	TLB-012 p.9; TLB-013 p.6
4	talabat pro 20-28% frequency uplift / 26-32% retention uplift (GCC+Jordan cohort, six markets pre-Dec 2024, excludes Egypt/Iraq)	ASM-016 (Approved)	TLB-001 p.18; TLB-015 pp.78,104; TLB-019 p.11
4	Rewards >15%, PostPaid 14% order-frequency uplift (Group)	Business_Relationships.md Chain 1	TLB-001, pp.18-19; TLB-026, p.134
4	AI/personalisation EBITDA contribution USD14mn+ (FY2024) → USD30mn+ (FY2025), Group	Topics/AI.md; Value_Driver_Tree_v2.md N-09	TLB-001 p.23; TLB-002 p.15
4	AI/personalisation EBITDA contribution not a discrete dollar line in either bucket	Value_Driver_Tree_v2.md N-06	—
4	Advertising revenue USD246mn→USD323mn (+32% y/y) vs. ~7%-of-GMV UAE benchmark not yet reached Group-wide	Topics/Revenue Drivers.md	TLB-001 p.21; TLB-002 pp.10,14,19; TLB-014 p.19
4	No disclosed ROI, payback period, or return figure for either investment bucket (Group, absence-of-disclosure finding)	Investment_Relationship_Map.md	corpus-wide
5	talabat's technology stack is licensed from Delivery Hero SE, a partner arrangement, not a standalone build	Entities/Technology_Platforms.md	Topics/AI.md
5	talabat pro live in 6 of 8 countries by end-2024 (excludes Egypt, Iraq)	Value_Driver_Tree_v2.md N-08/N-13	TLB-012, TLB-013, TLB-015, TLB-016, TLB-026
5	This plan's own decision-support system's MCP-integration status remains undecided	—	MEMORY.md
6	FY2024/FY2025 revenue by fee type (Commission, Delivery & Service, Subscription & Other Income, Advertising & Listing, Vouchers) — MECE by talabat's own Management Revenue convention	Strategic/Revenue Model.md; Facts/Revenue.md	TLB-001 p.27; TLB-002 p.20; TLB-009 p.4

14.3 The DEC-009 Allocation-Range Sensitivity Exhibit (Designated Home)

Action title: talabat's disclosed 175mn/120mn/55mn envelope is fixed; this OS's own five-option cost ranges within it are an analytical recommendation, not a talabat allocation.

Option	Bucket	Base	Upside	Downside	Basis
OPT-001 — talabat mart dark-store densification	Everyday App	USD 45-55mn	USD 55-70mn	USD 30-40mn	ASM-029
OPT-002 — talabat pro acceleration (Egypt, Iraq)	Everyday App (opex)	USD 8-12mn	USD 12-20mn	USD 3-6mn	ASM-030

Option	Bucket	Base	Upside	Downside	Basis
OPT-003 — Food-leadership CVP/partner-retention (UAE, Kuwait, Qatar)	Food-leadership	USD 25-32mn	USD 32-40mn	USD 15-22mn	ASM-031
OPT-004 — Advertising monetization gap closure	Enabling capability (cross-cutting, drawn from Everyday App opex)	USD 8-12mn	USD 12-18mn	USD 3-6mn	ASM-032
OPT-005 — AI/personalisation capability scaling	Enabling capability (cross-cutting, drawn from Everyday App opex)	USD 10-15mn	USD 15-22mn	USD 5-8mn	ASM-033

– Non-additive across the envelope. OPT-004 and OPT-005 are cross-cutting Enabling-capability options

14.4 Resumes of Key Team Members — Not Applicable

Stated honestly, not fabricated. Problem_Charter.md names six AASTMT MBA capstone team members (Group G02) with AASTMT student IDs, but every role assignment in that table is marked [assign] — no role has been finalized, and no biographical or professional-experience content for any team member exists anywhere in this project's corpus or vault.

14.5 Technical Specifications of the AI Product

Full detail is Section 5's — this appendix gives the compressed specification, keeping the same explicit separation Section 5 established between the two distinct AI systems in scope:

• talabat's own embedded AI/personalisation capability (scaled, not built, by the 2026 investment

14.6 Case Studies or Pilot Program Results

Two categories, kept explicitly separate — this section does not let either be mistaken for the other.

14.7 Legal and Compliance Documents

No separate legal or compliance filing was produced for, or by, this plan — this is a capstone Business Plan, not a regulatory submission, and no such document exists in this project's Outputs/ or elsewhere in the vault.

• Inter-company/data-governance mechanism: the Delivery Hero Group Inter-Company Data Transfer

Traceability (for this section's own claims)

Document Control and Validation

14/14 sections independently verified; all five whole-plan gates PASS. Human approval remains mandatory before capital movement.